

# COMPOSITIONAL ENTAILMENT LEARNING FOR HYPERBOLIC VISION-LANGUAGE MODELS

**Anonymous authors**

Paper under double-blind review

## ABSTRACT

Image-text representation learning forms a cornerstone in vision-language models, where pairs of images and textual descriptions are contrastively aligned in a shared embedding space. Since visual and textual concepts are naturally hierarchical, recent work has shown that hyperbolic space can serve as a high-potential manifold to learn vision-language representation with strong downstream performance. In this work, for the first time we show how to fully leverage the innate hierarchical nature of hyperbolic embeddings by looking beyond individual image-text pairs. We propose Compositional Entailment Learning for hyperbolic vision-language models. The idea is that an image is not only described by a sentence but is itself a composition of multiple object boxes, each with their own textual description. Such information can be obtained freely by extracting nouns from sentences and using openly available localized grounding models. We show how to hierarchically organize images, image boxes, and their textual descriptions through contrastive and entailment-based objectives. Empirical evaluation on a hyperbolic vision-language model trained with millions of image-text pairs shows that the proposed compositional learning approach outperforms conventional Euclidean CLIP learning, as well as recent hyperbolic alternatives, with better zero-shot and retrieval generalization and clearly stronger hierarchical performance. *Code to be released.*

## 1 INTRODUCTION

Vision-language modeling has witnessed rapid progress in recent years with innovative approaches such as CLIP (Radford et al., 2021) and ALIGN (Jia et al., 2021) using extensive vision-language data to train encoders for understanding visual and textual content simultaneously. Such encoders align visual scenes with textual descriptions in a shared high-dimensional Euclidean space, facilitating semantic understanding (Radford et al., 2021). While effective, conventional vision-language models only take a holistic approach to image-text representation learning, neglecting the intrinsic hierarchy and composition of elements within images. Indeed, a visual scene is commonly composed of multiple objects interacting with one another to form a precise context. See for example Fig. 1b with description: “*Mineral water with fresh herbs in a glass carafe on a garden table*”. Individually, these objects provide limited semantic meaning. Only through the interactions between these do we understand the specific context of both the scene and its parts, characterizing the single entities (cf. Fig. 1a). This object-scene hierarchy is analogous to a parent-child connection in a discrete tree where broader concepts are closer to the root while specific concepts reside deeper in the tree. These tree-like structures cannot be well represented in Euclidean space due its polynomial volume growth (Matoušek, 1999), whereas hyperbolic geometry does accommodate the exponential growth of trees (Gromov, 1987), making it more suitable for representing hierarchies.

Recently, Desai et al. (2023) introduced MERU, a hyperbolic contrastive vision-language model. MERU projects Euclidean embeddings from image and text encoders onto hyperbolic space and enforces *inter-modal* (text to image) partial ordering (Vendrov et al., 2016) using an entailment loss (Ganea et al., 2018a; Le et al., 2019) when optimizing encoder weights. Such hyperbolic image-text alignment has demonstrated strong quantitative performance on zero-shot downstream tasks, as well as increased interpretability of the shared embedding space. They, however, ignore the *intra-modal* hierarchical compositions of image-text pairs. Indeed, there is hierarchical semantics in language (Everaert et al., 2015), which has been leveraged to embed textual data in hyperbolic space (Dhingra et al., 2018). In the vision domain, work by Ge et al. (2023) uses object-centric scene

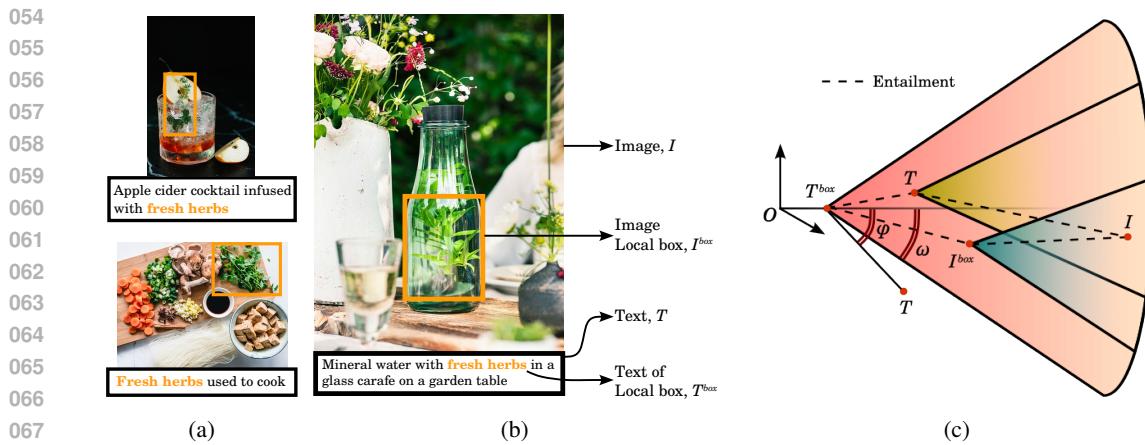


Figure 1: **Compositional Entailment Learning** for hyperbolic vision-language models. (a) same object appearing in different vision-language contexts (b) Visual-semantic ordering for an image-text pair:  $I$  (whole image) and  $T$  (full caption) provide context to the more general  $I^{box}$  (image local box) and  $T^{box}$  (text local box). (c) This specific-general ordering between  $(I, T)$ ,  $(I^{box}, T^{box})$ ,  $(I, I^{box})$ ,  $(T, T^{box})$  is enforced in hyperbolic space using entailment cones. The external angle  $\phi$  of a specific concept ( $T$ ) is pushed to be within the aperture threshold  $\eta\omega$  of the general concept ( $T^{box}$ ).

hierarchies to learn a hyperbolic space where visually similar objects are clustered near the origin and scenes consisting of them are descendants. Zhong et al. (2022) propose RegionCLIP that only learns regional representations using contrastive learning and knowledge distillation. These prior works beg the question of what strategy can be adapted to compound the individual benefits of the inter-modal hierarchy and the two intra-modal hierarchies to encompass *scene and region* level understanding.

To this end, we introduce Hyperbolic Compositional CLIP (HyCoCLIP), a contrastive learning method that accounts for compositional orders in both inter-modal and intra-modal settings in hyperbolic space. We approach the problem by using explicit hierarchies while training the encoders. This hierarchy is constituted of object crops (*image boxes*) within an image and corresponding nouns/phrases (*text boxes*) within the text as broader concepts of the whole image-text concept. We outline a robust hierarchical learning objective by using both entire images and image boxes, as well as complete captions and text boxes. This strategy involves both inter-modal hierarchies, where text generally provides broader context than images, and intra-modal hierarchies, where we consider the “boxes” more general than the complete image. In the hierarchical spatial representation, broader concepts are embedded near the origin of the metric space, while more fine-grained concepts are positioned towards the border, akin to tree graphs, see Fig. 1c.

We show that HyCoCLIP outperforms CLIP and MERU on zero-shot image classification and is competitive on zero-shot retrieval and object detection when trained on a 20M pre-training dataset. Additionally, we show that HyCoCLIP improves on hierarchical classification tasks compared to the baselines and that its representation space is more interpretable and hierarchically aligned. Our contributions are summarized as follows: (1) We introduce HyCoCLIP for learning vision-language representations in a shared hyperbolic space using scene compositions that are semantically and hierarchically aligned. (2) We propose Compositional Entailment Learning, where image-text compositions are optimized through hyperbolic contrastive and entailment cone losses. (3) We demonstrate empirically that HyCoCLIP is more hierarchically aware and is highly competitive to existing vision-language models.

## 2 HYPERBOLIC COMPOSITIONAL CLIP - HYCoCLIP

We propose a compositional learning scheme that enforces the semantic alignment of latent representations in the hyperbolic space, explicitly modeling intra- and inter-modal relationships of visual and language data by leveraging their joint hierarchical nature. Here, we first provide a short background

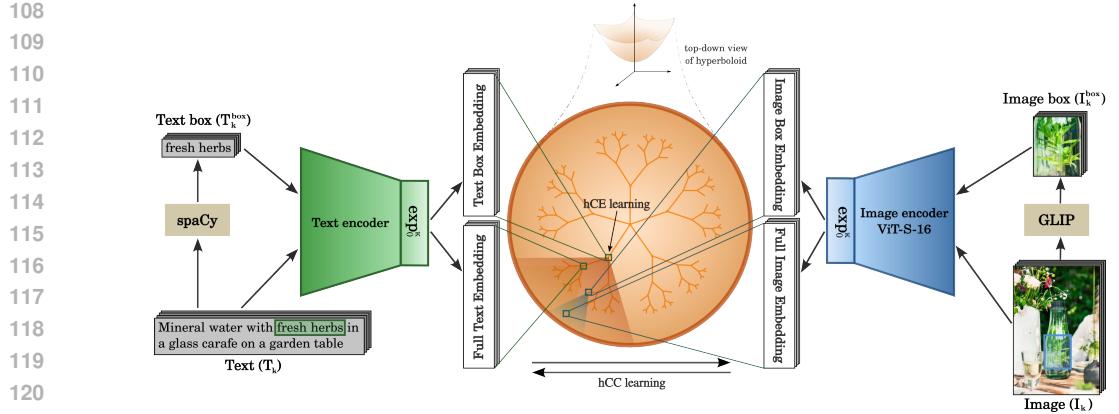


Figure 2: **An overview of HyCoCLIP.** Text and image boxes are extracted offline from image-text datasets (sides). Next, HyCoCLIP’s encoder modules embed the images and texts, projecting the representations in the hyperbolic latent space. HyCoCLIP preserves the inter-modal and intra-modal relationships by accommodating broader/finer concepts close to the center/border and by using entailment cones to give an interpretable structure to the learned latent space (cf. Fig. 1c).

with the required hyperbolic functions to make such compositional learning possible. Afterward, we outline our compositional encoding of image-text pairs.

## 2.1 BACKGROUND

Hyperbolic geometry is a non-Euclidean geometry characterized by a constant negative curvature. The resulting space has the desirable property that volumes of subsets can grow exponentially as a function of their radius, making it an ideal choice for learning representations of data with an inherent hierarchical or tree-like structure (Sarkar, 2011; Nickel & Kiela, 2017; Krioukov et al., 2010). While several isometric models are used in literature for modeling hyperbolic space, we limit our background discussion to the Lorentz (or hyperboloid) model used in this work and refer to Cannon et al. (1997); Peng et al. (2022) for detailed information on the other models.

The Lorentz model, denoted by  $\mathbb{L}^n$ , is an  $n$ -dimensional manifold represented as the upper sheet of a two-sheeted hyperboloid in  $(n+1)$ -dimensional Minkowski spacetime. For each vector  $\mathbf{p} \in \mathbb{R}^{n+1}$ , the first dimension is taken as the *time*-axis, denoted  $p_0$ , and the remaining  $n$  dimensions as the *spatial*-coordinates, denoted  $\tilde{\mathbf{p}} \in \mathbb{R}^n$ . This model is described as

$$\mathbb{L}^n = \left\{ \mathbf{p} \in \mathbb{R}^{n+1} : \langle \mathbf{p}, \mathbf{p} \rangle_{\mathbb{L}} = -\frac{1}{\kappa}, p_0 = \sqrt{1/\kappa + \|\tilde{\mathbf{p}}\|^2}, \kappa > 0 \right\}, \quad (1)$$

where  $-\kappa \in \mathbb{R}$  is the curvature of the space and  $\langle \cdot, \cdot \rangle_{\mathbb{L}}$  is the Lorentzian inner product defined for  $\mathbf{p}, \mathbf{q} \in \mathbb{L}^n$  as

$$\langle \mathbf{p}, \mathbf{q} \rangle_{\mathbb{L}} = -p_0 q_0 + \langle \tilde{\mathbf{p}}, \tilde{\mathbf{q}} \rangle_{\mathbb{E}}, \quad (2)$$

with  $\langle \cdot, \cdot \rangle_{\mathbb{E}}$  denoting the Euclidean inner product. The Lorentzian distance between two points in  $\mathbb{L}^n$  is the length of the shortest path (*geodesic*) connecting them, which can be computed as

$$d_{\mathbb{L}}(\mathbf{p}, \mathbf{q}) = \sqrt{1/\kappa} \cdot \cosh^{-1}(-\kappa \langle \mathbf{p}, \mathbf{q} \rangle_{\mathbb{L}}), \quad \mathbf{p}, \mathbf{q} \in \mathbb{L}^n. \quad (3)$$

This metric induces the Lorentzian norm  $\|\mathbf{p}\|_{\mathbb{L}} = \langle \mathbf{p}, \mathbf{p} \rangle_{\mathbb{L}}$ . The tangent space  $T_{\mathbf{p}}\mathbb{L}^n$  is well-defined for all the points  $\mathbf{p} \in \mathbb{L}^n$ , and the exponential map represents the projecting map from the tangent space to the hyperboloid. Given a point  $\mathbf{v} \in T_{\mathbf{p}}\mathbb{L}^n$  the exponential map can be computed as

$$\exp_{\mathbf{p}}^{\kappa}(\mathbf{v}) = \cosh(\sqrt{\kappa}\|\mathbf{v}\|_{\mathbb{L}})\mathbf{p} + \frac{\sinh(\sqrt{\kappa}\|\mathbf{v}\|_{\mathbb{L}})}{\sqrt{\kappa}\|\mathbf{v}\|_{\mathbb{L}}} \mathbf{v}. \quad (4)$$

Such a map can be used to move from Euclidean space to hyperbolic space by considering Euclidean vectors to be tangent vectors at the origin  $\mathbf{0} = (\sqrt{1/\kappa}, 0, \dots, 0)^T$  of the hyperbolic space and using  $\exp_{\mathbf{0}}^{\kappa}$  to project these onto the hyperboloid (Khrulkov et al., 2020).

162 2.2 COMPOSITIONAL ENTAILMENT LEARNING  
163

164 We strive to learn the hierarchical compositional relations of images, boxes, and textual descriptions.  
 165 Our idea is based on the following observation: the content inside a box of an image is hierarchically  
 166 more general than the entire image. While counter-intuitive at first glance, Fig. 1b shows why this  
 167 is the case: the box shows an object and the entire image additionally shows the context in which  
 168 the object occurs, making it a semantically more specific scenario. From a hyperbolic perspective,  
 169 semantically general/broad concepts are embedded closer to the origin, while more fine-grained  
 170 concepts are positioned towards the border, akin to tree graphs (cf. Fig. 1c).  
 171

172 In this work, we are given a dataset  $D = \{(I_k, T_k)\}_{k=1}^K$  of  $K$  image-text pairs. Our goal is to train  
 173 image and text encoders with a shared embedding space to align the visual and semantic inputs. The  
 174 method is summarized in Fig. 2. Let  $(I_k^{\text{box}}, T_k^{\text{box}})$  be the local box with a short description from an  
 175 image-text pair obtained following the automated procedure detailed in Appendix A. We propose  
 176 a Compositional Entailment Learning objective in hyperbolic space to optimize the hierarchical  
 177 compositions. Our approach consists of two parts, namely a compositional contrastive loss and a  
 178 compositional entailment loss which we discuss next.  
 179

180 **Hierarchical Compositional Contrastive (hCC) learning.** Image-text models commonly rely  
 181 on contrastive objectives to align and distribute the multi-modal data. In our approach, we rely on  
 182 hyperbolic embeddings to align visuals and text. Let  $f_I(\cdot)$  and  $f_T(\cdot)$  denote arbitrary encoders for the  
 183 image and text inputs respectively. And, let  $g_I(I_k) = \exp_0^\kappa(f_I(I_k))$  and  $g_T(T_k) = \exp_0^\kappa(f_T(T_k))$   
 184 denote the hyperbolic representation of image  $I_k$  and textual description  $T_k$  respectively. To compute  
 185 the contrastive loss over image-text pairs in a batch  $B$ , we take the negative Lorentzian distance as  
 186 our similarity metric and formulate it with the softmax, using temperature  $\tau$ , for a batch of size  $(B)$   
 187 containing images ( $I$ ) and text ( $T$ ) as  
 188

$$L_{\text{cont}}^*(I, T) = - \sum_{i \in B} \log \frac{\exp(d_{\mathbb{L}}(g_I(I_i), g_T(T_i))/\tau)}{\sum_{k=1, k \neq i}^B \exp(d_{\mathbb{L}}(g_I(I_i), g_T(T_k))/\tau)}, \quad (5)$$

189 where negatives for an image are picked from the texts in the batch. Similarly, we can define the  
 190 loss when picking negatives for a text from images in the batch as  $L_{\text{cont}}^*(T, I)$ . To extend such a  
 191 contrastive setup with our image-text compositions, we have to consider that due to the generalized  
 192 information in a box, different images in a batch can have similar box-level information. To avoid  
 193 unwanted negatives in a batch, we only contrast the box image with other entire images, and vice  
 194 versa which have specific information. This avoids negative alignment between image-box pairs  
 195 and boxes from different images. The final hierarchical Compositional Contrastive (hCC) loss is  
 196 formulated as  
 197

$$hCC(I, T, I^{\text{box}}, T^{\text{box}}) = \frac{1}{4} \left( \underbrace{L_{\text{cont}}^*(I, T) + L_{\text{cont}}^*(T, I)}_{\text{specific-info contrast}} + \underbrace{L_{\text{cont}}^*(I^{\text{box}}, T) + L_{\text{cont}}^*(T^{\text{box}}, I)}_{\text{general-info contrast}} \right). \quad (6)$$

200 **Hierarchical Compositional Entailment (hCE) learning.** Ganea et al. (2018a) introduced hyper-  
 201 bolic entailment cones that generalize the idea of partial order embeddings (Vendrov et al., 2016) by  
 202 using the inherent hierarchical structure of the hyperbolic space. Entailment cones define a region  
 203  $\mathfrak{R}_q$  for every possible point  $q$  in the space such that all points  $p \in \mathfrak{R}_q$  are semantically linked to  
 204  $q$  as its child concepts. As such, points in  $\mathfrak{R}_q$  are expected to contain specific information for the  
 205 general concept  $q$ . Considering the Lorentz model  $\mathbb{L}^n$ , the half-aperture of these conical regions ( $\mathfrak{R}_q$ )  
 206 is formulated by Le et al. (2019); Desai et al. (2023) as  
 207

$$\omega(q) = \sin^{-1} \left( \frac{2K}{\sqrt{\kappa} \|\tilde{q}\|} \right), \quad (7)$$

208 where  $-\kappa$  is the curvature of the space and a constant  $K = 0.1$  is set to limit values near the origin  
 209 (see Ganea et al. (2018a)). The aperture inversely depends on the norm  $\|\tilde{q}\|$ . Inferring from this, a  
 210 general concept with a wider aperture would lie closer to the origin. A specific concept would have a  
 211 narrower aperture and lie further from the origin in the hyperbolic space.  
 212

213 To learn partial orders in this space, specific concepts  $p$  must be pushed to be within the aperture  $\omega(q)$ .  
 214 This is done by penalizing encoders with the angular residual of outward point  $p$  having an exterior  
 215 angle  $\phi(p, q)$  as shown in Fig. 1c. This is formulated by Le et al. (2019); Desai et al. (2023) as  
 216

216

$$L_{ent}(\mathbf{p}, \mathbf{q}) = \max(0, \phi(\mathbf{p}, \mathbf{q}) - \omega(\mathbf{q})), \quad (8)$$

217 where the exterior angle is given by,

$$\phi(\mathbf{p}, \mathbf{q}) = \cos^{-1} \left( \frac{p_0 + q_0 \kappa \langle \mathbf{p}, \mathbf{q} \rangle_{\mathbb{L}}}{\|\tilde{\mathbf{q}}\| \sqrt{(\kappa \langle \mathbf{p}, \mathbf{q} \rangle_{\mathbb{L}})^2 - 1}} \right). \quad (9)$$

220  
221  
222  
223  
224  
225  
226  
227  
228 However, intuitively observing the entailment loss presented in Eq. 8 shows that this loss would push any outward point  $\mathbf{p}$  only towards the  $\mathfrak{R}_q$  region’s border. Hence, we add a threshold to the half-aperture  $\omega(\mathbf{q})$  effectively making it flexible to accommodate  $\mathbf{p}$  at various spatial distances from  $\mathbf{q}$ , see Fig. 3. We reformulate Eq. 8 with half-aperture threshold  $\eta$  as

$$L_{ent}^*(\mathbf{p}, \mathbf{q}) = \max(0, \phi(\mathbf{p}, \mathbf{q}) - \eta \omega(\mathbf{q})). \quad (10)$$

229  
230  
231 Entailment cones enable us to enforce the hierarchical image-text relations given by the compositions. We formulate the Hierarchical Compositional Entailment (hCE) loss by considering that images and textual descriptions are not identical, but that text precedes image, akin to Desai et al. (2023). We additionally consider the relation *whole  $\Rightarrow$  box* for both images and texts. Hence, the hCE loss would comprise both image-text inter-modality entailments and text-text, image-image intra-modality entailments as

$$hCE(I, T, I^{box}, T^{box}) = \underbrace{L_{ent}^*(I^{box}, T^{box}) + L_{ent}^*(I, T)}_{\text{inter-modality entailment}} + \underbrace{L_{ent}^*(I, I^{box}) + L_{ent}^*(T, T^{box})}_{\text{intra-modality entailment}}. \quad (11)$$

232  
233 In Fig. 1c we visualize how the image-text compositions should be organized in hyperbolic compositional entailment.

234  
235 **Hierarchical Compositional (hC) learning.** We aggregate the losses to form the overall hierarchical Compositional (hC) loss for HyCoCLIP by taking a weighted sum of the two loss components:

$$hC = hCC + \gamma hCE. \quad (12)$$

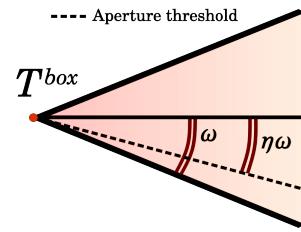
236  
237 In Appendix B, we detail all hyperparameters, thresholds, and further implementation details.

238  
239  
240 **Computational complexity.** Our approach enables us to double the amount of visual and textual data to learn from. The training time scales linearly with the increase in training volume; for ViT-B/16, HyCoCLIP requires 73 hours of training, compared to 46 hours for MERU and 45 hours for CLIP. We note that our method inference maintains the same efficiency as CLIP and MERU and allows for scalable deployment.

### 255 3 EXPERIMENTS

#### 256 3.1 BENCHMARK

257  
258 **Datasets** We develop our models using grounded vision-language pairs. This could be human-annotated such as the Localized narratives subset of Open Images (Pont-Tuset et al., 2020) or the Flickr30K Entities dataset (Plummer et al., 2015). However, the sizes of these datasets are fairly limited considering the intensive efforts of manual labelling. Hence, we depend on automatic grounded information generated by pre-trained phrase grounding models. Several large-scale grounded language-vision datasets are publicly available by Li et al. (2023) and Peng et al. (2023). We train our models using the large-scale training corpus - Grounded Image-Text Pairs (GRIT) dataset (Peng et al., 2023) containing 20.5 million grounded vision-language pairs which are processed from the even larger COYO-700M (Byeon et al., 2022) dataset. Information on the grounding procedure is added in Appendix A. We similarly use the grounding procedure on the RedCaps dataset (Desai et al., 2021) originally used to train MERU. Additionally, we use the smaller-scale grounded Conceptual Captions 3M (CC3M) (Li et al., 2023; Sharma et al., 2018) dataset for hyperparameter search.



259  
260 **Figure 3: Aperture threshold**  
261  $\eta$  scaling the aperture  $\omega$  to increase or decrease the width of  
262 the entailment cone.

Table 1: **Zero-shot image classification evaluation.**  $\dagger$  denotes reproduced results from MERU. When using boxes during pre-training, numbers in squared brackets represent the additional box-pairs counts. For RedCaps, we find results for CLIP and MERU consistent with Desai et al. (2023) even when trained with a smaller batch size. **Bold-face** numbers are the best performances **overall**, while **underlined** figures are the best within all the models sharing the same ViT backbone. Our method outperforms baselines on 15 out of the 16 evaluation datasets.

	w/ boxes	samples (M)	General datasets						Fine-grained datasets						Misc. datasets				
			ImageNet	CIFAR-10	CIFAR-100	SUN397	Caltech-101	STL-10	Food-101	CUB	Cars	Aircraft	Pets	Flowers	DTD	EuroSAT	RESISC45		
																	Country211		
<b>RedCaps</b>																			
ViT S/16	CLIP $\dagger$	$\times$	11.4	32.5	66.7	35.8	26.7	60.8	89.8	<u>72.5</u>	29.8	<u>11.1</u>	1.3	72.5	44.9	16.4	30.1	27.7	5.0
	CLIP	$\checkmark$	11.4 [6.3]	<u>30.2</u>	76.5	<u>42.4</u>	25.8	62.3	89.5	<u>69.6</u>	<u>25.7</u>	8.5	<u>2.2</u>	<u>65.3</u>	38.6	13.6	<u>36.6</u>	28.5	4.6
	MERU $\dagger$	$\times$	11.4	31.4	65.9	35.2	26.8	58.1	89.3	71.4	29.0	8.3	1.6	71.0	40.9	17.0	29.9	29.3	4.7
	MERU	$\checkmark$	11.4 [6.3]	29.9	76.4	39.9	26.6	62.3	89.5	68.4	25.4	8.9	1.2	67.2	37.6	13.0	30.5	27.6	4.3
	HyCoCLIP	$\checkmark$	5.8 [6.3]	<u>31.9</u>	<u>77.4</u>	<u>37.7</u>	<u>27.6</u>	<u>64.5</u>	<u>90.9</u>	<u>71.1</u>	<u>28.8</u>	<u>9.7</u>	<u>1.1</u>	<u>70.5</u>	<u>41.4</u>	<u>13.4</u>	<u>22.7</u>	<u>30.7</u>	<u>4.4</u>
<b>GRIT</b>																			
ViT S/16	CLIP	$\times$	20.5	36.7	70.2	42.6	49.5	73.6	89.7	44.7	9.8	6.9	2.0	44.6	14.8	22.3	<b>40.7</b>	40.1	5.1
	CLIP	$\checkmark$	20.5 [35.9]	36.2	84.2	<u>54.8</u>	46.1	74.1	91.6	43.2	11.9	6.0	2.5	45.9	18.1	24.0	32.4	35.5	4.7
	MERU	$\times$	20.5	35.4	71.2	42.0	48.6	73.0	89.8	48.8	10.9	6.5	2.3	42.7	17.3	18.6	39.1	38.9	5.3
	MERU	$\checkmark$	20.5 [35.9]	35.0	85.0	54.0	44.6	73.9	91.6	41.1	10.1	5.6	2.2	43.9	15.9	24.5	39.3	33.5	4.8
	HyCoCLIP	$\checkmark$	20.5 [35.9]	<u>41.7</u>	<u>85.0</u>	53.6	<u>52.5</u>	<u>75.7</u>	<u>92.5</u>	<u>50.2</u>	<u>14.7</u>	<u>8.1</u>	<u>4.2</u>	<u>52.0</u>	<u>20.5</u>	<u>22.3</u>	33.8	<b>45.7</b>	5.2
ViT B/16	CLIP	$\times$	20.5	40.6	78.9	48.3	53.0	76.7	92.4	48.6	10.0	9.0	3.4	45.9	21.3	23.4	37.1	42.7	5.7
	MERU	$\times$	20.5	40.1	78.6	49.3	53.0	72.8	93.2	51.5	11.9	8.6	3.7	48.5	21.2	22.2	<u>31.7</u>	44.2	5.6
	HyCoCLIP	$\checkmark$	20.5 [35.9]	<b>45.8</b>	<b>88.8</b>	<b>60.1</b>	<b>57.2</b>	<b>81.3</b>	<b>95.0</b>	<b>59.2</b>	<b>16.4</b>	<b>11.6</b>	<b>3.7</b>	<b>56.8</b>	<b>23.9</b>	<b>29.4</b>	35.8	45.6	<b>6.5</b>

**Baseline Comparisons** We compare HyCoCLIP against CLIP and MERU. We reproduce the CLIP and MERU models by training on the RedCaps dataset, reducing the batch size to 768 to fit on our available compute. We further retrain CLIP and MERU from scratch on the GRIT dataset. To fairly evaluate the impact of including image-text boxes, we also retrain CLIP and MERU when image-text boxes are included as additional pre-training samples.

### 3.2 DOWNSTREAM TASKS

To assess performance, we evaluate HyCoCLIP on several downstream tasks. For zero-shot image classification, the label set is fitted to multiple prompts which are embedded using the text encoder and then averaged to obtain a single embedding per label. The closest text embedding is picked from the collection as the predicted class for an image. We report the model’s accuracy on 16 image classification datasets. Similarly, we assess our method on zero-shot retrieval tasks to determine if complex concepts, like scenes and captions, are accurately preserved in the representation space. Further, we evaluate the models on object detection task to analyze the regional understanding of HyCoCLIP. We also evaluate the hierarchical nature of HyCoCLIP using multiple hierarchical metrics. Additionally, we assess the scene understanding capability of HyCoCLIP on two compositional benchmarks - VL-Checklist (Zhao et al., 2022) and VG Attribution (Yüksekgönül et al., 2023).

**Zero-shot image classification** From Table 1, we find our reproduced results for CLIP and MERU are fairly consistent with Desai et al. (2023) even when trained with smaller batch size on RedCaps. On grounding RedCaps and filtering noise, we notice only 5.8 million image-text pairs are retained containing 6.3 million boxes. Training the baselines with these additional boxes also demonstrates reduced performance. Alternatively, the quantity of data is significantly higher for GRIT with 20.5 million image-text pairs and 35.9 million boxes in total. To differentiate between the datasets, we compare the ratio of the box area with the image area of all data points and plot a histogram in Fig. 4. A lower ratio signifies that phrases have more localized information in the image and constitute a better semantic parent for the whole image. This is evident for GRIT while boxes generated for RedCaps do not seem to localize well. We, therefore, recommend GRIT over RedCaps for grounded pre-training and thus report the results of models pre-trained on GRIT. We find that HyCoCLIP

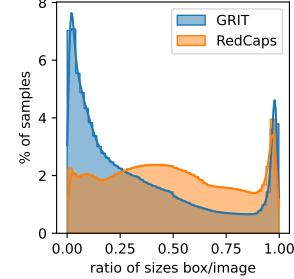


Figure 4: **Histogram of ratios** of box area wrt the full image for GRIT and RedCaps. The latter reports generally larger crops, indicating lower precision in grounding concepts.

324  
 325 **Table 2: Zero-shot retrieval, detection, and hierarchical classification.** HyCoCLIP performs best  
 326 in image retrieval, and hierarchical classification and is competitive in text retrieval. **Bold** figures  
 327 indicate the best results for each ViT backbone.

Vision encoder	Model	w/ boxes	Text retrieval				Image retrieval				Hierarchical metrics				
			COCO		Flickr		COCO		Flickr						
			R@5	R@10	R@5	R@10	R@5	R@10	R@5	R@10	TIE( $\downarrow$ )	LCA( $\downarrow$ )	J( $\uparrow$ )	P <sub>H</sub> ( $\uparrow$ )	R <sub>H</sub> ( $\uparrow$ )
ViT S/16	CLIP	$\times$	69.3	79.1	<b>90.2</b>	<b>95.2</b>	53.7	65.2	81.1	87.9	4.02	2.39	0.76	0.83	0.84
	CLIP	$\checkmark$	60.7	71.8	84.2	91.3	47.1	58.6	73.1	82.1	4.03	2.38	0.76	0.83	0.83
	MERU	$\times$	68.8	78.8	89.4	94.8	53.6	65.3	80.4	87.5	4.08	2.39	0.76	0.83	0.83
	MERU	$\checkmark$	<b>72.7</b>	<b>81.9</b>	83.5	90.1	46.6	58.3	60.0	71.7	4.08	2.39	0.75	0.83	0.83
	HyCoCLIP	$\checkmark$	69.5	79.5	89.1	93.9	<b>55.2</b>	<b>66.6</b>	<b>81.5</b>	<b>88.1</b>	<b>3.55</b>	<b>2.17</b>	<b>0.79</b>	<b>0.86</b>	<b>0.85</b>
	CLIP	$\times$	71.4	81.5	<b>93.6</b>	<b>96.9</b>	57.4	68.5	83.5	89.9	3.60	2.21	0.79	0.85	0.85
B/16	MERU	$\times$	<b>72.3</b>	82.0	93.5	96.2	57.4	68.6	84.0	90.0	3.63	2.22	0.78	0.85	0.85
	HyCoCLIP	$\checkmark$	72.0	<b>82.0</b>	92.6	95.4	<b>58.4</b>	<b>69.3</b>	<b>84.9</b>	<b>90.3</b>	<b>3.17</b>	<b>2.05</b>	<b>0.81</b>	<b>0.87</b>	<b>0.87</b>

336  
 337 performs best across a wide range of datasets and settings when pre-training is done on GRIT. We  
 338 especially note the performance on ImageNet, where we obtain an accuracy of 45.8% compared  
 339 to 40.1% (MERU) and 40.6% (CLIP). Interestingly, adding image-text boxes to CLIP and MERU  
 340 training does not improve performance, despite nearly doubling the training samples.  
 341

342 **Zero-shot retrieval** For the retrieval task, the top-k image/text embeddings are picked from a  
 343 collection for input text/image embedding based on the distance score (Eq. 3). We perform this  
 344 task zero-shot on the COCO validation set (Lin et al., 2014) and the Flickr30K test set (Young et al.,  
 345 2014; Karpathy & Fei-Fei, 2015). We show the retrieval results in Table 2. We find that our method  
 346 performs slightly worse on Flickr text retrieval while demonstrating increased performance on image  
 347 retrieval over CLIP and MERU. We also note a significant decrease in the performance of CLIP and  
 348 MERU when adding local information. These results further highlight the need for our approach.  
 349 Naively adding these boxes as additional samples is not effective because the boxes are often without  
 350 broader context, and the text is highly generic compared to the whole images. Only by optimizing for  
 351 their hierarchical compositional nature as done in our approach is it possible to get better performance.  
 352 Our method aims to obtain a hierarchically aligned representation space, but this is not necessarily  
 353 beneficial for the task of retrieval, where proximity of text and image embeddings is key. Regardless,  
 354 our approach remains highly competitive.  
 355

356 **Hierarchical Classification** A characteristic feature of hyperbolic spaces is their ability to represent  
 357 hierarchical structures present in data. We evaluate our models for this property on several hierarchical  
 358 classification metrics (Kosmopoulos et al., 2015) described in Appendix C. We use the WordNet  
 359 hierarchy (Miller, 1994) of the ImageNet class labels (Deng et al., 2009; Russakovsky et al., 2015) for  
 360 the hierarchical classification task. The image classification setup is kept similar and the final scores  
 361 are averaged over the validation set. Table 2 reports the results of HyCoCLIP and other baselines on  
 362 these metrics. We observe a consistent improvement, confirming that the hierarchy of the class labels  
 363 is better represented in its embedding space.  
 364

365 **Zero-shot object detection** We utilize pre-trained vision-  
 366 language models to recognize proposed object regions.  
 367 Specifically, we evaluate the scenario where ground-truth  
 368 bounding boxes from the COCO detection dataset are used  
 369 as region proposals and predict the correct categories with  
 370 a setup similar to image classification. We compare our  
 371 method with RegionCLIP (Zhong et al., 2022) whose vision  
 372 encoder (ResNet50x4) was trained on CC3M with a frozen  
 373 text encoder (originally trained on CLIP400M). We report  
 374 the average precision (AP) on the 17 novel categories data  
 375 split (Bansal et al., 2018). As shown in Table 3, HyCoCLIP  
 376 outperforms the baselines, surpassing RegionCLIP on the  
 377 novel categories. We believe this highlights a key advantage  
 378 of our approach—its ability to leverage inherent hierarchies  
 379 for more effective semantic concept alignment.

380  
 381 **Table 3: Zero-shot object detection** with ground-truth boxes evaluated  
 382 on COCO 17 novel categories split (Bansal et al., 2018). HyCoCLIP  
 383 shows the best average precision (AP).

Model	AP
CLIP	51.2
MERU	55.8
RegionCLIP	65.2
HyCoCLIP	<b>68.5</b>

378  
 379 **Table 4: Ablation study on loss terms during**  
 380 **pre-training** of HyCoCLIP-ViT-S/16 on grounded  
 381 CC3M evaluated for image classification and  
 382 Flickr image/text retrieval. Lower accuracy/R@5  
 383 indicates a more influential loss term.

	classification		retrieval		
	ImageNet	Food-101	Mean(16)	Text	Image
<b>Pre-training losses</b>					
HyCoCLIP	16.7	10.6	22.3	56.2	46.4
<b>hCC loss</b>					
$-L_{cont}^*(I, T)$	16.0	9.4	22.5	55.2	46.2
$-L_{cont}^*(T, I)$	16.1	10.2	22.2	55.2	45.4
$-L_{cont}^*(I^{box}, T)$	13.8	8.7	19.3	49.1	42.9
$-L_{cont}^*(T^{box}, I)$	15.2	7.6	20.4	55.9	44.5
<b>hCE loss</b>					
$-L_{ent}^*(I, T)$	14.9	9.9	21.8	54.3	46.1
$-L_{ent}^*(I^{box}, T^{box})$	16.1	9.3	22.3	54.8	45.4
$-L_{ent}^*(I, I^{box})$	16.1	10.1	21.5	55.0	45.6
$-L_{ent}^*(T, T^{box})$	16.3	9.2	22.3	55.6	45.1

383  
 384 **Table 5: Scene understanding evaluation.**  
 385 HyCoCLIP show better performance on both  
 386 benchmarks indicating good object comprehen-  
 387 sion of the visual scene.

Model	VL-CO	VG-A
CLIP	49.3	63.3
MERU	50.5	61.8
HyCoCLIP	<b>59.8</b>	<b>68.4</b>

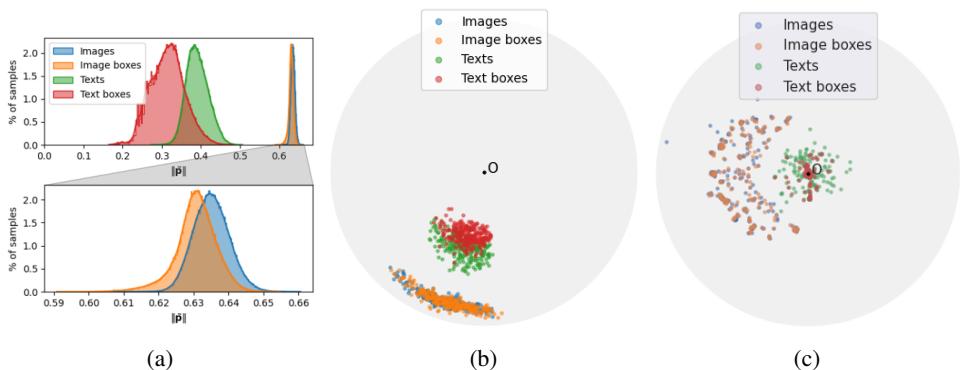
388  
 389 **Table 6: Ablation study on batch size** shows  
 390 saturation after 768.

Batch size	ImageNet
512	11.1
640	11.3
768	12.2
896	12.2
1024	12.1
1536	12.1

391 **Scene Understanding** Given the compositional pre-training strategy, we expect HyCoCLIP to  
 392 be provisioned with localized object/noun information in both vision and language and improve  
 393 upon such aspects. The setup for these benchmarks is the same, for a given image the model has to  
 394 pick between the correct caption and a hard negative caption. For more information on the methods  
 395 used to generate the hard negative captions, we refer to Appendix E. From Table 5, we see that  
 396 CLIP and MERU give near-random performance for the VL-Checklist-Object (VL-CO) (Zhao et al.,  
 397 2022) benchmark in which object information in the captions is perturbed. HyCoCLIP improves  
 398 considerably on these experiments reaching 60% accuracy indicating good object comprehension  
 399 of the visual scene. HyCoCLIP also performs well on VG-Attribution (VG-A) (Yüksekgnl et al.,  
 400 2023) reporting a mean accuracy of 68.4% surpassing other methods. We refer to Appendix E for  
 401 further analysis.

### 402 3.3 ABLATION STUDY

403 **Pre-training loss terms** We examine the impact of the terms in hCC (Eq. 6) and hCE (Eq. 11)  
 404 losses by pre-training the model several times, each time turning off a single loss term. We use the  
 405 grounded CC3M dataset and train for 40k steps. Table 4 shows the results of this experiment. A lower  
 406 accuracy and recall on image classification and retrieval respectively, indicate a higher influence of  
 407 corresponding loss term. For hCC loss, we find that our hypothesis of contrasting the generalized



408 **Figure 5: Visualizing the learned hyperbolic space of HyCoCLIP in lower dimensions** using  
 409 samples from GRIT. (a) distribution of embedding distances from the origin, HyCoCLIP embeds  
 410 text data closer to the origin wrt the images and boxes samples with a smaller radius wrt full  
 411 images/captions. On the right, (b) HoroPCA and (c) CO-SNE visualizations of the latent space in  $\mathbb{L}^2$ .

432 information in boxes against entire images and text is indeed beneficial. For hCE loss, we see that the  
 433 terms entailing the image ( $I$ ) are most influential.  
 434

435 **Scaling w.r.t batch size** We train our models using a batch size 768 according to available compute  
 436 (Appendix B). To study the influence of this hyperparameter, we train our primary baseline MERU-  
 437 ViT-S using CC3M for various batch sizes and report their zero-shot performance on ImageNet  
 438 classification. Table 6 indicates no empirical benefits when working with larger batch sizes. In  
 439 the contrastive setting, the number of positives grows linearly, and the number of negatives grows  
 440 quadratically in a batch. When using softmax, the ratio of positives to negatives affects loss functions  
 441 differently depending on the type of similarity metric that is being used. This can explain the  
 442 difference in batch size behavior of our approach. The saturation of softmax loss with increasing  
 443 batch size has been previously discussed by Zhai et al. (2023), and the entailment loss may also  
 444 contribute to this early saturation.

## 4 ANALYZING THE HYPERBOLIC SPACE

447 **Visualizing the learned hyperbolic space** We visualize  
 448 the learned hyperbolic space in lower dimensions to see if  
 449 the image, text, and corresponding box embeddings are dis-  
 450 tributed in a proper semantic hierarchy. To this end, we plot  
 451 the distribution of the spatial norms of 128k random sam-  
 452 ples of training data in a histogram. Furthermore, we use  
 453 *HoroPCA* (Chami et al., 2021) for reducing the dimension-  
 454 ality for 200 image-text pairs along with their boxes. Lastly,  
 455 we extract 50 principal components to suppress noise and use  
 456 CO-SNE (Guo et al., 2022) to bring the embeddings to the  
 457 low-dimensional space.

458 Fig. 5a shows that the embedding distributions of texts and  
 459 their corresponding boxes are well separated, while images  
 460 and their box representations display similar norms. This  
 461 spatial contraction in image embeddings arises from the  
 462 convergence of contrastive loss within a confined entailment  
 463 cone, as noted by Ramasinghe et al. (2024). Furthermore,  
 464 many image boxes are almost identical to the full image (cf.  
 465 Fig. 4), making it challenging for the network to differentiate  
 466 between them. Nonetheless, the bottom plot in Fig. 5a shows  
 467 that the box embeddings distribute closer to the origin, thus  
 468 displaying hierarchical ordering. From Fig. 5b and 5c, we  
 469 observe the semantic separation in the two principal compo-  
 470 nents of HoroPCA and in the 2D space formed with CO-SNE,  
 471 indicating an apparent hierarchy between the components.

472 **Interpolating between points in hyperbolic space** We in-  
 473 terpolate the geodesic connecting an image (source) with an-  
 474 other image (target) and also with the origin similar to Desai  
 475 et al. (2023), which have been visualized on the bottom-right  
 476 of Fig. 6. This intuitively represents traversing between nodes  
 477 in a discrete tree. This is useful in visualizing the ancestors  
 478 of any given image and qualitatively verifying the hierarchi-  
 479 cal properties of the learned hyperbolic space. We do the  
 480 shortest path traversal in the tangent space details of which  
 481 are in Appendix G. We use grounded Flickr30K (Li et al.,  
 482 2023) to generate the collection of representations of images,  
 483 texts, and corresponding boxes. Fig. 6 shows the result of  
 484 100 points being interpolated between two randomly selected images from pexels.com as well  
 485 as to the origin. We observe that HyCoCLIP can fetch representations of both data modes in a very  
 rational hierarchy. More interpolation examples are added in Appendix I where we also compare  
 interpolation in MERU and CLIP representation space.

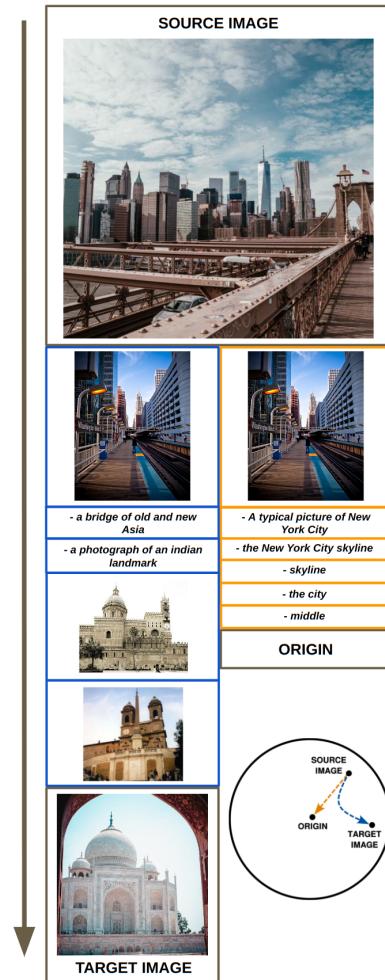


Figure 6: **Interpolation between points.** Multimodal retrieval results when moving from (top) an image to (left) another image or (right) the origin, as depicted in the (bottom-right) circle.

486    **5 RELATED WORK**

487

488    **Vision-language models** Currently expanding at a rapid pace, this topic has been in focus for multi-  
 489    ple decades with initial works in image retrieval (Mori et al., 1999), semantic segmentation (Barnard  
 490    et al., 2003), and object classification (Wang et al., 2009) leveraging natural language descriptors  
 491    for computer vision. Later works (He & Peng, 2017) utilized more expressive representations from  
 492    multi-modal neural network encoders. The advent of transformers (Vaswani et al., 2017) and vision  
 493    transformers (Dosovitskiy et al., 2021) helped construct a highly semantic embedding space for  
 494    texts and images, respectively. Recent works have explored creating a shared embedding space by  
 495    leveraging various pre-training strategies to integrate text and image information. We refer the reader  
 496    to the survey by Gan et al. (2022) for a comprehensive overview. Many approaches use contrastive  
 497    learning as a core method, like CLIP (Radford et al., 2021) and ALIGN (Jia et al., 2021). Zhao  
 498    et al. (2022) advanced this with RegionCLIP, which aligns image regions with textual concepts for  
 499    object detection. More recently, MERU (Desai et al., 2023) combines entailment learning (Ganea  
 500    et al., 2018a; Le et al., 2019) with the CLIP approach to learn embeddings in hyperbolic space  
 501    capturing latent visual-semantic hierarchies. We extend this to include image patches and caption  
 502    parts, enforcing an ordering that reflects the hierarchy shared by both modalities.

503

504    **Learning in hyperbolic space** Hyperbolic space for representation learning has desirable properties  
 505    for data with an inherent hierarchical or tree-like structure (Nickel & Kiela, 2017; Chamberlain et al.,  
 506    2017). When generating embeddings in hyperbolic space from such data, its innate hierarchical  
 507    structure can be retained with minimal distortion. As a result, hyperbolic deep learning has rapidly  
 508    gained traction (Peng et al., 2022; Mettes et al., 2023). Recent works have developed methods  
 509    for building neural networks that operate in hyperbolic space (Ganea et al., 2018b; Shimizu et al.,  
 510    2021) and corresponding optimization algorithms (Bécigneul & Ganea, 2019; Bonnabel, 2013). This  
 511    led to the use of hyperbolic models in many different modalities such as graphs (Liu et al., 2019),  
 512    text (Dhingra et al., 2018; Tifrea et al., 2019), images (van Spengler et al., 2023; Atigh et al., 2022),  
 513    videos (Long et al., 2020), etc. Other recent work has focused on combining embedding spaces of  
 514    different modalities (Liu et al., 2020; Desai et al., 2023). Our work similarly learns multimodal  
 515    representations in hyperbolic space to benefit from its inductive hierarchical bias.

516    **Hierarchies in vision and language** Vendrov et al. (2016) use a visual-semantic hierarchy over  
 517    words, sentences, and images to learn representations in a supervised fashion. They consider  
 518    hypernym-hyponym relations in language to construct a hierarchy. This concept has been used in  
 519    hypernymy detection tasks (Nguyen et al., 2017; Vulic & Mrksic, 2018). Hierarchies formed by  
 520    constituency-based parse trees have been used to learn embeddings in hyperbolic space by Dhingra  
 521    et al. (2018). In vision, several works sought to connect scenes to objects and parts of objects within  
 522    the scene. Early works have used such information for pose estimation, image segmentation, and  
 523    object and contour detection (Bourdev & Malik, 2009; Arbeláez et al., 2011). Recently, un-/self-  
 524    supervised methods have been used for representation learning leveraging hierarchical segmentation  
 525    of an image by Zhang & Maire (2020) and object-scene hypernymy by Xie et al. (2021); Ge et al.  
 526    (2023). We combine hypernymy relations of vision and language.

527

528    **6 CONCLUSION**

529

530    The idea of this work is to use object compositions within a scene and its description, along with  
 531    the visual-semantic ordering between image and text to learn hyperbolic representations that are  
 532    semantically and hierarchically aligned. Our proposed HyCoCLIP improves over standard CLIP  
 533    and its recent hyperbolic extension MERU in zero-shot classification. Moreover, our approach has  
 534    increased scene understanding and better hierarchical structuring. Further, we qualitatively analyze  
 535    the space by visualizing representations and through point-to-point interpolation which substantiates  
 536    HyCoCLIP’s ability to embed multi-modal hierarchies in a shared space. The method has certain  
 537    limitations, with a key challenge being the need to generate bounding box information from image-  
 538    caption pairs during training. This increases the volume of visual and textual data processed by  
 539    HyCoCLIP, though it still preserves scalability during inference. Additionally, while our hierarchical  
 540    training strategy improves interpretability by separating images and texts into distinct regions in the  
 541    embedding space, it may not be optimal for tasks like large-scale retrieval.

540 REFERENCES  
541

- 542 Morris Alper and Hadar Averbuch-Elor. Emergent visual-semantic hierarchies in image-text  
543 representations. *CoRR*, abs/2407.08521, 2024. doi: 10.48550/ARXIV.2407.08521. URL  
544 <https://doi.org/10.48550/arXiv.2407.08521>.
- 545 Pablo Arbeláez, Michael Maire, Charless Fowlkes, and Jitendra Malik. Contour detection and  
546 hierarchical image segmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*,  
547 33(5):898–916, 2011. doi: 10.1109/TPAMI.2010.161.
- 548 Mina Ghadimi Atigh, Julian Schoep, Erman Acar, Nanne van Noord, and Pascal Mettes. Hyperbolic  
549 image segmentation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*,  
550 *CVPR 2022, New Orleans, LA, USA, June 18-24, 2022*, pp. 4443–4452. IEEE, 2022. doi: 10.  
551 1109/CVPR52688.2022.00441. URL <https://doi.org/10.1109/CVPR52688.2022.00441>.
- 552 Ankan Bansal, Karan Sikka, Gaurav Sharma, Rama Chellappa, and Ajay Divakaran. Zero-shot  
553 object detection. In Vittorio Ferrari, Martial Hebert, Cristian Sminchisescu, and Yair Weiss (eds.),  
554 *Computer Vision - ECCV 2018 - 15th European Conference, Munich, Germany, September 8-14,  
555 2018, Proceedings, Part I*, volume 11205 of *Lecture Notes in Computer Science*, pp. 397–414.  
556 Springer, 2018. doi: 10.1007/978-3-030-01246-5\_24. URL [https://doi.org/10.1007/978-3-030-01246-5\\_24](https://doi.org/10.1007/978-3-030-01246-5_24).
- 557 Kobus Barnard, Pinar Duygulu, David A. Forsyth, Nando de Freitas, David M. Blei, and Michael I.  
558 Jordan. Matching words and pictures. *J. Mach. Learn. Res.*, 3:1107–1135, 2003. URL <http://jmlr.org/papers/v3/barnard03a.html>.
- 559 Gary Bécigneul and Octavian-Eugen Ganea. Riemannian adaptive optimization methods. In *7th  
560 International Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA,  
561 May 6-9, 2019*. OpenReview.net, 2019. URL <https://openreview.net/forum?id=r1eiqi09K7>.
- 562 Silvère Bonnabel. Stochastic gradient descent on riemannian manifolds. *IEEE Trans. Autom.  
563 Control.*, 58(9):2217–2229, 2013. doi: 10.1109/TAC.2013.2254619. URL <https://doi.org/10.1109/TAC.2013.2254619>.
- 564 Lubomir Bourdev and Jitendra Malik. Poselets: Body part detectors trained using 3d human pose  
565 annotations. In *2009 IEEE 12th International Conference on Computer Vision*, pp. 1365–1372,  
566 2009. doi: 10.1109/ICCV.2009.5459303.
- 567 Minwoo Byeon, Beomhee Park, Haecheon Kim, Sungjun Lee, Woonhyuk Baek, and Sae-  
568 hoon Kim. Coyo-700m: Image-text pair dataset. [https://github.com/kakaobrain/  
569 coyo-dataset](https://github.com/kakaobrain/coyo-dataset), 2022.
- 570 James W Cannon, William J Floyd, Richard Kenyon, Walter R Parry, et al. Hyperbolic geometry.  
571 *Flavors of geometry*, 31(59-115):2, 1997.
- 572 Benjamin Paul Chamberlain, James R. Clough, and Marc Peter Deisenroth. Neural embeddings of  
573 graphs in hyperbolic space. *CoRR*, abs/1705.10359, 2017. URL <http://arxiv.org/abs/1705.10359>.
- 574 Ines Chami, Albert Gu, Dat Nguyen, and Christopher Ré. Horopca: Hyperbolic dimensionality  
575 reduction via horospherical projections. In Marina Meila and Tong Zhang (eds.), *Proceedings of  
576 the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual  
577 Event*, volume 139 of *Proceedings of Machine Learning Research*, pp. 1419–1429. PMLR, 2021.  
578 URL <http://proceedings.mlr.press/v139/chami21a.html>.
- 579 Xinlei Chen, Saining Xie, and Kaiming He. An empirical study of training self-supervised vision  
580 transformers. In *2021 IEEE/CVF International Conference on Computer Vision, ICCV 2021, Mon-  
581 treal, QC, Canada, October 10-17, 2021*, pp. 9620–9629. IEEE, 2021. doi: 10.1109/ICCV48922.  
582 2021.00950. URL <https://doi.org/10.1109/ICCV48922.2021.00950>.

- 594 Ofer Dekel, Joseph Keshet, and Yoram Singer. Large margin hierarchical classification. In Carla E.  
 595 Brodley (ed.), *Machine Learning, Proceedings of the Twenty-first International Conference (ICML*  
 596 *2004), Banff, Alberta, Canada, July 4-8, 2004*, volume 69 of *ACM International Conference*  
 597 *Proceeding Series*. ACM, 2004. doi: 10.1145/1015330.1015374. URL <https://doi.org/10.1145/1015330.1015374>.
- 598  
 599 Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale  
 600 hierarchical image database. In *2009 IEEE Computer Society Conference on Computer Vision and*  
 601 *Pattern Recognition (CVPR 2009), 20-25 June 2009, Miami, Florida, USA*, pp. 248–255. IEEE  
 602 Computer Society, 2009. doi: 10.1109/CVPR.2009.5206848. URL <https://doi.org/10.1109/CVPR.2009.5206848>.
- 603  
 604 Karan Desai, Gaurav Kaul, Zubin Aysola, and Justin Johnson. Redcaps: Web-curated image-  
 605 text data created by the people, for the people. In Joaquin Vanschoren and Sai-Kit Yeung  
 606 (eds.), *Proceedings of the Neural Information Processing Systems Track on Datasets and*  
 607 *Benchmarks 1, NeurIPS Datasets and Benchmarks 2021, December 2021, virtual*, 2021. URL  
 608 <https://datasets-benchmarks-proceedings.neurips.cc/paper/2021/hash/e00da03b685a0dd18fb6a08af0923de0-Abstract-round1.html>.
- 609  
 610 Karan Desai, Maximilian Nickel, Tanmay Rajpurohit, Justin Johnson, and Shanmukha Ramakr-  
 611 ishna Vedantam. Hyperbolic image-text representations. In Andreas Krause, Emma Brunskill,  
 612 Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett (eds.), *International*  
 613 *Conference on Machine Learning, ICML 2023, 23-29 July 2023, Honolulu, Hawaii, USA*, vol-  
 614 *ume 202 of Proceedings of Machine Learning Research*, pp. 7694–7731. PMLR, 2023. URL  
 615 <https://proceedings.mlr.press/v202/desai23a.html>.
- 616  
 617 Bhuwan Dhingra, Christopher J. Shallue, Mohammad Norouzi, Andrew M. Dai, and George E. Dahl.  
 618 Embedding text in hyperbolic spaces. In Goran Glavas, Swapna Somasundaran, Martin Riedl, and  
 619 Eduard H. Hovy (eds.), *Proceedings of the Twelfth Workshop on Graph-Based Methods for Natural*  
 620 *Language Processing, TextGraphs@NAACL-HLT 2018, New Orleans, Louisiana, USA, June 6,*  
 621 *2018*, pp. 59–69. Association for Computational Linguistics, 2018. doi: 10.18653/V1/W18-1708.  
 622 URL <https://doi.org/10.18653/v1/w18-1708>.
- 623  
 624 Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas  
 625 Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit,  
 626 and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale.  
 627 In *9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria,*  
 628 *May 3-7, 2021*. OpenReview.net, 2021. URL <https://openreview.net/forum?id=YicbFdNTTy>.
- 629  
 630 Martin B.H. Everaert, Marinus A.C. Huybregts, Noam Chomsky, Robert C. Berwick, and Johan J.  
 631 Bolhuis. Structures, not strings: Linguistics as part of the cognitive sciences. *Trends in Cog-*  
 632 *nitive Sciences*, 19(12):729–743, 2015. ISSN 1364-6613. doi: <https://doi.org/10.1016/j.tics.2015.09.008>. URL <https://www.sciencedirect.com/science/article/pii/S1364661315002326>.
- 633  
 634 Zhe Gan, Linjie Li, Chunyuan Li, Lijuan Wang, Zicheng Liu, and Jianfeng Gao. Vision-language  
 635 pre-training: Basics, recent advances, and future trends. *Found. Trends Comput. Graph. Vis.*,  
 636 14(3-4):163–352, 2022. doi: 10.1561/0600000105. URL <https://doi.org/10.1561/0600000105>.
- 637  
 638 Octavian-Eugen Ganea, Gary Bécigneul, and Thomas Hofmann. Hyperbolic entailment cones for  
 639 learning hierarchical embeddings. In Jennifer G. Dy and Andreas Krause (eds.), *Proceedings of the*  
 640 *35th International Conference on Machine Learning, ICML 2018, Stockholmsmässan, Stockholm,*  
 641 *Sweden, July 10-15, 2018*, volume 80 of *Proceedings of Machine Learning Research*, pp. 1632–  
 642 1641. PMLR, 2018a. URL <http://proceedings.mlr.press/v80/ganea18a.html>.
- 643  
 644 Octavian-Eugen Ganea, Gary Bécigneul, and Thomas Hofmann. Hyperbolic neural networks. In  
 645 Samy Bengio, Hanna M. Wallach, Hugo Larochelle, Kristen Grauman, Nicolò Cesa-Bianchi, and  
 646 Roman Garnett (eds.), *Advances in Neural Information Processing Systems 31: Annual Conference*  
 647 *on Neural Information Processing Systems 2018, NeurIPS 2018, December 3-8, 2018, Montréal,*

- 648      Canada, pp. 5350–5360, 2018b. URL <https://proceedings.neurips.cc/paper/2018/hash/dbab2adc8f9d078009ee3fa810bea142-Abstract.html>.
- 649  
650  
651      Songwei Ge, Shlok Mishra, Simon Kornblith, Chun-Liang Li, and David Jacobs. Hyperbolic  
652      contrastive learning for visual representations beyond objects. In *IEEE/CVF Conference on*  
653      *Computer Vision and Pattern Recognition, CVPR 2023, Vancouver, BC, Canada, June 17-24, 2023*,  
654      pp. 6840–6849. IEEE, 2023. doi: 10.1109/CVPR52729.2023.00661. URL <https://doi.org/10.1109/CVPR52729.2023.00661>.
- 655  
656      Mikhael Gromov. Hyperbolic groups. In *Essays in group theory*, pp. 75–263. Springer, 1987.
- 657  
658      Yunhui Guo, Haoran Guo, and Stella X. Yu. CO-SNE: dimensionality reduction and visualization for  
659      hyperbolic data. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR*  
660      *2022, New Orleans, LA, USA, June 18-24, 2022*, pp. 11–20. IEEE, 2022. doi: 10.1109/CVPR52688.  
661      2022.00011. URL <https://doi.org/10.1109/CVPR52688.2022.00011>.
- 662      Xiangteng He and Yuxin Peng. Fine-grained image classification via combining vision and language.  
663      In *2017 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2017, Honolulu, HI, USA, July 21-26, 2017*, pp. 7332–7340. IEEE Computer Society, 2017. doi: 10.1109/CVPR.2017.775.  
664      URL <https://doi.org/10.1109/CVPR.2017.775>.
- 665      Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. spacy: Industrial-  
666      strength natural language processing in python. 2020. doi: 10.5281/zenodo.1212303.
- 667  
668      Chao Jia, Yinfei Yang, Ye Xia, Yi-Ting Chen, Zarana Parekh, Hieu Pham, Quoc V. Le, Yun-Hsuan  
669      Sung, Zhen Li, and Tom Duerig. Scaling up visual and vision-language representation learning  
670      with noisy text supervision. In Marina Meila and Tong Zhang (eds.), *Proceedings of the 38th*  
671      *International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event*,  
672      volume 139 of *Proceedings of Machine Learning Research*, pp. 4904–4916. PMLR, 2021. URL  
673      <http://proceedings.mlr.press/v139/jia21b.html>.
- 674  
675      Andrej Karpathy and Li Fei-Fei. Deep visual-semantic alignments for generating image descriptions.  
676      In *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2015, Boston, MA,*  
677      *USA, June 7-12, 2015*, pp. 3128–3137. IEEE Computer Society, 2015. doi: 10.1109/CVPR.2015.  
678      7298932. URL <https://doi.org/10.1109/CVPR.2015.7298932>.
- 679      Valentin Khrulkov, Leyla Mirvakhabova, Evgeniya Ustinova, Ivan Oseledets, and Victor Lempitsky.  
680      Hyperbolic image embeddings. In *Proceedings of the IEEE/CVF conference on computer vision*  
681      *and pattern recognition*, pp. 6418–6428, 2020.
- 682      Aris Kosmopoulos, Ioannis Partalas, Éric Gaussier, Georgios Palioras, and Ion Androutsopoulos.  
683      Evaluation measures for hierarchical classification: a unified view and novel approaches. *Data*  
684      *Min. Knowl. Discov.*, 29(3):820–865, 2015. doi: 10.1007/S10618-014-0382-X. URL <https://doi.org/10.1007/s10618-014-0382-x>.
- 685  
686      Dmitri Krioukov, Fragkiskos Papadopoulos, Maksim Kitsak, Amin Vahdat, and Marián Boguñá.  
687      Hyperbolic geometry of complex networks. *Phys. Rev. E*, 82:036106, Sep 2010. doi: 10.1103/  
688      PhysRevE.82.036106. URL <https://link.aps.org/doi/10.1103/PhysRevE.82.036106>.
- 689  
690      Matt Le, Stephen Roller, Laetitia Papaxanthos, Douwe Kiela, and Maximilian Nickel. Inferring  
691      concept hierarchies from text corpora via hyperbolic embeddings. In Anna Korhonen, David R.  
692      Traum, and Lluís Márquez (eds.), *Proceedings of the 57th Conference of the Association for*  
693      *Computational Linguistics, ACL 2019, Florence, Italy, July 28- August 2, 2019, Volume 1: Long*  
694      *Papers*, pp. 3231–3241. Association for Computational Linguistics, 2019. doi: 10.18653/V1/  
695      P19-1313. URL <https://doi.org/10.18653/v1/p19-1313>.
- 696  
697      Liunian Harold Li, Pengchuan Zhang, Haotian Zhang, Jianwei Yang, Chunyuan Li, Yiwu Zhong, Li-  
698      juan Wang, Lu Yuan, Lei Zhang, Jenq-Neng Hwang, Kai-Wei Chang, and Jianfeng Gao. Grounded  
699      language-image pre-training. In *IEEE/CVF Conference on Computer Vision and Pattern Recog-  
700      nition, CVPR 2022, New Orleans, LA, USA, June 18-24, 2022*, pp. 10955–10965. IEEE, 2022.  
701      doi: 10.1109/CVPR52688.2022.01069. URL <https://doi.org/10.1109/CVPR52688.2022.01069>.

- 702 Yuheng Li, Haotian Liu, Qingyang Wu, Fangzhou Mu, Jianwei Yang, Jianfeng Gao, Chunyuan  
 703 Li, and Yong Jae Lee. GLIGEN: open-set grounded text-to-image generation. In *IEEE/CVF*  
 704 *Conference on Computer Vision and Pattern Recognition, CVPR 2023, Vancouver, BC, Canada,*  
 705 *June 17-24, 2023*, pp. 22511–22521. IEEE, 2023. doi: 10.1109/CVPR52729.2023.02156. URL  
 706 <https://doi.org/10.1109/CVPR52729.2023.02156>.
- 707 Tsung-Yi Lin, Michael Maire, Serge J. Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr  
 708 Dollár, and C. Lawrence Zitnick. Microsoft COCO: common objects in context. In David J.  
 709 Fleet, Tomás Pajdla, Bernt Schiele, and Tinne Tuytelaars (eds.), *Computer Vision - ECCV 2014*  
 710 - 13th European Conference, Zurich, Switzerland, September 6-12, 2014, Proceedings, Part V,
- 711 volume 8693 of *Lecture Notes in Computer Science*, pp. 740–755. Springer, 2014. doi: 10.1007/  
 712 978-3-319-10602-1\48. URL [https://doi.org/10.1007/978-3-319-10602-1\\_48](https://doi.org/10.1007/978-3-319-10602-1_48).
- 713 Qi Liu, Maximilian Nickel, and Douwe Kiela. Hyperbolic graph neural networks. In Hanna M.  
 714 Wallach, Hugo Larochelle, Alina Beygelzimer, Florence d’Alché-Buc, Emily B. Fox, and Roman  
 715 Garnett (eds.), *Advances in Neural Information Processing Systems 32: Annual Conference on*  
 716 *Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver,*  
 717 *BC, Canada*, pp. 8228–8239, 2019. URL <https://proceedings.neurips.cc/paper/2019/hash/103303dd56a731e377d01f6a37badae3-Abstract.html>.
- 718 Shaoteng Liu, Jingjing Chen, Liangming Pan, Chong-Wah Ngo, Tat-Seng Chua, and Yu-Gang Jiang.  
 719 Hyperbolic visual embedding learning for zero-shot recognition. In *2020 IEEE/CVF Conference*  
 720 *on Computer Vision and Pattern Recognition, CVPR 2020, Seattle, WA, USA, June 13-19, 2020*, pp.  
 721 9270–9278. Computer Vision Foundation / IEEE, 2020. doi: 10.1109/CVPR42600.2020.00929.  
 722 URL [https://openaccess.thecvf.com/content\\_CVPR\\_2020/html/Liu\\_Hyperbolic\\_Visual\\_EMBEDDING\\_Learning\\_for\\_Zero-Shot\\_Recognition\\_CVPR\\_2020\\_paper.html](https://openaccess.thecvf.com/content_CVPR_2020/html/Liu_Hyperbolic_Visual_EMBEDDING_Learning_for_Zero-Shot_Recognition_CVPR_2020_paper.html).
- 723 Teng Long, Pascal Mettes, Heng Tao Shen, and Cees G. M. Snoek. Searching for actions on the  
 724 hyperbole. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June  
 725 2020.
- 726 Ilya Loshchilov and Frank Hutter. SGDR: stochastic gradient descent with warm restarts. In *5th*  
 727 *International Conference on Learning Representations, ICLR 2017, Toulon, France, April 24-26,*  
 728 *2017, Conference Track Proceedings*. OpenReview.net, 2017. URL <https://openreview.net/forum?id=Skq89Scxx>.
- 729 Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *7th International*  
 730 *Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA, May 6-9, 2019*.  
 731 OpenReview.net, 2019. URL <https://openreview.net/forum?id=Bkg6RiCqY7>.
- 732 Jiří Matoušek. On embedding trees into uniformly convex banach spaces. *Israel Journal of Mathematics*, 114(1):221–237, 1999.
- 733 Pascal Mettes, Mina Ghadimi Atigh, Martin Keller-Ressel, Jeffrey Gu, and Serena Yeung. Hyperbolic  
 734 deep learning in computer vision: A survey. *CoRR*, abs/2305.06611, 2023. doi: 10.48550/ARXIV.  
 735 2305.06611. URL <https://doi.org/10.48550/arXiv.2305.06611>.
- 736 George A. Miller. WordNet: A lexical database for English. In *Human Language Technology:*  
 737 *Proceedings of a Workshop held at Plainsboro, New Jersey, March 8-11, 1994*, 1994. URL  
 738 <https://aclanthology.org/H94-1111>.
- 739 Yasuhide Mori, Hironobu Takahashi, and Ryu ichi Oka. Image-to-word transformation based on divid-  
 740 ing and vector quantizing images with words. 1999. URL <https://api.semanticscholar.org/CorpusID:18574318>.
- 741 Kim Anh Nguyen, Maximilian Köper, Sabine Schulte im Walde, and Ngoc Thang Vu. Hierarchical  
 742 embeddings for hypernymy detection and directionality. In Martha Palmer, Rebecca Hwa, and  
 743 Sebastian Riedel (eds.), *Proceedings of the 2017 Conference on Empirical Methods in Natural*  
 744 *Language Processing, EMNLP 2017, Copenhagen, Denmark, September 9-11, 2017*, pp. 233–  
 745 243. Association for Computational Linguistics, 2017. doi: 10.18653/V1/D17-1022. URL  
 746 <https://doi.org/10.18653/v1/d17-1022>.

- 756 Maximilian Nickel and Douwe Kiela. Poincaré embeddings for learning hierarchical representations.  
 757 In Isabelle Guyon, Ulrike von Luxburg, Samy Bengio, Hanna M. Wallach, Rob Fergus, S. V. N.  
 758 Vishwanathan, and Roman Garnett (eds.), *Advances in Neural Information Processing Systems  
 759 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017,  
 760 Long Beach, CA, USA*, pp. 6338–6347, 2017. URL <https://proceedings.neurips.cc/paper/2017/hash/59dfa2df42d9e3d41f5b02bfc32229dd-Abstract.html>.
- 761
- 762 Wei Peng, Tuomas Varanka, Abdelrahman Mostafa, Henglin Shi, and Guoying Zhao. Hyperbolic  
 763 deep neural networks: A survey. *IEEE Trans. Pattern Anal. Mach. Intell.*, 44(12):10023–10044,  
 764 2022. doi: 10.1109/TPAMI.2021.3136921. URL <https://doi.org/10.1109/TPAMI.2021.3136921>.
- 765
- 766 Zhiliang Peng, Wenhui Wang, Li Dong, Yaru Hao, Shaohan Huang, Shuming Ma, and Furu Wei.  
 767 Kosmos-2: Grounding multimodal large language models to the world. *CoRR*, abs/2306.14824,  
 768 2023. doi: 10.48550/ARXIV.2306.14824. URL <https://doi.org/10.48550/arXiv.2306.14824>.
- 769
- 770 Bryan A. Plummer, Liwei Wang, Chris M. Cervantes, Juan C. Caicedo, Julia Hockenmaier, and  
 771 Svetlana Lazebnik. Flickr30k entities: Collecting region-to-phrase correspondences for richer  
 772 image-to-sentence models. In *2015 IEEE International Conference on Computer Vision, ICCV  
 773 2015, Santiago, Chile, December 7-13, 2015*, pp. 2641–2649. IEEE Computer Society, 2015. doi:  
 774 10.1109/ICCV.2015.303. URL <https://doi.org/10.1109/ICCV.2015.303>.
- 775
- 776 Jordi Pont-Tuset, Jasper R. R. Uijlings, Soravit Changpinyo, Radu Soricut, and Vittorio Ferrari.  
 777 Connecting vision and language with localized narratives. In Andrea Vedaldi, Horst Bischof,  
 778 Thomas Brox, and Jan-Michael Frahm (eds.), *Computer Vision - ECCV 2020 - 16th European  
 779 Conference, Glasgow, UK, August 23-28, 2020, Proceedings, Part V*, volume 12350 of *Lecture  
 780 Notes in Computer Science*, pp. 647–664. Springer, 2020. doi: 10.1007/978-3-030-58558-7\\_38.  
 781 URL [https://doi.org/10.1007/978-3-030-58558-7\\_38](https://doi.org/10.1007/978-3-030-58558-7_38).
- 782
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal,  
 783 Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever.  
 784 Learning transferable visual models from natural language supervision. In Marina Meila and  
 785 Tong Zhang (eds.), *Proceedings of the 38th International Conference on Machine Learning, ICML  
 786 2021, 18-24 July 2021, Virtual Event*, volume 139 of *Proceedings of Machine Learning  
 787 Research*, pp. 8748–8763. PMLR, 2021. URL <http://proceedings.mlr.press/v139/radford21a.html>.
- 788
- 789 Sameera Ramasinghe, Violetta Shevchenko, Gil Avraham, and Ajanthan Thalaiyasingam. Accept the  
 790 modality gap: An exploration in the hyperbolic space. In *Proceedings of the IEEE/CVF Conference  
 791 on Computer Vision and Pattern Recognition (CVPR)*, pp. 27263–27272, June 2024.
- 792
- Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang,  
 793 Andrej Karpathy, Aditya Khosla, Michael S. Bernstein, Alexander C. Berg, and Li Fei-Fei. Imagenet  
 794 large scale visual recognition challenge. *Int. J. Comput. Vis.*, 115(3):211–252, 2015. doi: 10.  
 795 1007/S11263-015-0816-Y. URL <https://doi.org/10.1007/s11263-015-0816-y>.
- 796
- Rik Sarkar. Low distortion delaunay embedding of trees in hyperbolic plane. In Marc J. van Kreveld  
 797 and Bettina Speckmann (eds.), *Graph Drawing - 19th International Symposium, GD 2011, Eind-  
 798 hoven, The Netherlands, September 21-23, 2011, Revised Selected Papers*, volume 7034 of *Lecture  
 799 Notes in Computer Science*, pp. 355–366. Springer, 2011. doi: 10.1007/978-3-642-25878-7\\_34.  
 800 URL [https://doi.org/10.1007/978-3-642-25878-7\\_34](https://doi.org/10.1007/978-3-642-25878-7_34).
- 801
- Piyush Sharma, Nan Ding, Sebastian Goodman, and Radu Soricut. Conceptual captions: A cleaned,  
 802 hypernymed, image alt-text dataset for automatic image captioning. In Iryna Gurevych and Yusuke  
 803 Miyao (eds.), *Proceedings of the 56th Annual Meeting of the Association for Computational  
 804 Linguistics, ACL 2018, Melbourne, Australia, July 15-20, 2018, Volume 1: Long Papers*, pp.  
 805 2556–2565. Association for Computational Linguistics, 2018. doi: 10.18653/V1/P18-1238. URL  
 806 <https://aclanthology.org/P18-1238/>.
- 807
- Ryohei Shimizu, Yusuke Mukuta, and Tatsuya Harada. Hyperbolic neural networks++. In *9th Interna-  
 808 tional Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3-7, 2021.  
 809 OpenReview.net*, 2021. URL <https://openreview.net/forum?id=Ec85b0tUwbA>.

- 810 Alexandru Tifrea, Gary Bécigneul, and Octavian-Eugen Ganea. Poincare glove: Hyperbolic word  
 811 embeddings. In *7th International Conference on Learning Representations, ICLR 2019, New*  
 812 *Orleans, LA, USA, May 6-9, 2019*. OpenReview.net, 2019. URL <https://openreview.net/forum?id=Ske5r3AqK7>.
- 813
- 814 Hugo Touvron, Matthieu Cord, Matthijs Douze, Francisco Massa, Alexandre Sablayrolles, and Hervé  
 815 Jégou. Training data-efficient image transformers & distillation through attention. In Marina  
 816 Meila and Tong Zhang (eds.), *Proceedings of the 38th International Conference on Machine*  
 817 *Learning, ICML 2021, 18-24 July 2021, Virtual Event*, volume 139 of *Proceedings of Machine*  
 818 *Learning Research*, pp. 10347–10357. PMLR, 2021. URL <http://proceedings.mlr.press/v139/touvron21a.html>.
- 819
- 820
- 821 Max van Spengler, Erwin Berkhouit, and Pascal Mettes. Poincaré resnet. In *IEEE/CVF International*  
 822 *Conference on Computer Vision, ICCV 2023, Paris, France, October 1-6, 2023*, pp. 5396–5405.  
 823 IEEE, 2023. doi: 10.1109/ICCV51070.2023.00499. URL <https://doi.org/10.1109/ICCV51070.2023.00499>.
- 824
- 825
- 826 Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez,  
 827 Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In Isabelle Guyon, Ulrike von  
 828 Luxburg, Samy Bengio, Hanna M. Wallach, Rob Fergus, S. V. N. Vishwanathan, and Roman  
 829 Garnett (eds.), *Advances in Neural Information Processing Systems 30: Annual Conference on*  
 830 *Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*, pp.  
 831 5998–6008, 2017. URL <https://proceedings.neurips.cc/paper/2017/hash/3f5ee243547dee91fdbd053c1c4a845aa-Abstract.html>.
- 832
- 833 Ivan Vendrov, Ryan Kiros, Sanja Fidler, and Raquel Urtasun. Order-embeddings of images and  
 834 language. In Yoshua Bengio and Yann LeCun (eds.), *4th International Conference on Learning*  
 835 *Representations, ICLR 2016, San Juan, Puerto Rico, May 2-4, 2016, Conference Track Proceedings*,  
 836 2016. URL <http://arxiv.org/abs/1511.06361>.
- 837
- 838 Ivan Vulic and Nikola Mrksic. Specialising word vectors for lexical entailment. In Marilyn A. Walker,  
 839 Heng Ji, and Amanda Stent (eds.), *Proceedings of the 2018 Conference of the North American*  
 840 *Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-*  
 841 *HLT 2018, New Orleans, Louisiana, USA, June 1-6, 2018, Volume 1 (Long Papers)*, pp. 1134–  
 842 1145. Association for Computational Linguistics, 2018. doi: 10.18653/V1/N18-1103. URL  
 843 <https://doi.org/10.18653/v1/n18-1103>.
- 844 Josiah Wang, Katja Markert, and Mark Everingham. Learning models for object recognition from  
 845 natural language descriptions. In Andrea Cavallaro, Simon Prince, and Daniel C. Alexander  
 846 (eds.), *British Machine Vision Conference, BMVC 2009, London, UK, September 7-10, 2009.*  
 847 *Proceedings*, pp. 1–11. British Machine Vision Association, 2009. doi: 10.5244/C.23.2. URL  
 848 <https://doi.org/10.5244/C.23.2>.
- 849 Jiahao Xie, Xiaohang Zhan, Ziwei Liu, Yew Soon Ong, and Chen Change Loy. Unsu-  
 850 pervised object-level representation learning from scene images. In Marc’Aurelio Ran-  
 851 zato, Alina Beygelzimer, Yann N. Dauphin, Percy Liang, and Jennifer Wortman Vaughan  
 852 (eds.), *Advances in Neural Information Processing Systems 34: Annual Conference on Neu-*  
 853 *ral Information Processing Systems 2021, NeurIPS 2021, December 6-14, 2021, virtual*, pp.  
 854 28864–28876, 2021. URL <https://proceedings.neurips.cc/paper/2021/hash/f1b6f2857fb6d44dd73c7041e0aa0f19-Abstract.html>.
- 855
- 856 Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual  
 857 denotations: New similarity metrics for semantic inference over event descriptions. *Trans. Assoc.*  
 858 *Comput. Linguistics*, 2:67–78, 2014. doi: 10.1162/TACL\_A\_00166. URL [https://doi.org/10.1162/tacl\\_a\\_00166](https://doi.org/10.1162/tacl_a_00166).
- 859
- 860 Mert Yüksekgönül, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and  
 861 why vision-language models behave like bags-of-words, and what to do about it? In *The Eleventh*  
 862 *International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5,*  
 863 2023. OpenReview.net, 2023. URL <https://openreview.net/pdf?id=KRLUvxh8uaX>.

- 864 Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language  
 865 image pre-training. In *IEEE/CVF International Conference on Computer Vision, ICCV 2023, Paris,  
 866 France, October 1-6, 2023*, pp. 11941–11952. IEEE, 2023. doi: 10.1109/ICCV51070.2023.01100.  
 867 URL <https://doi.org/10.1109/ICCV51070.2023.01100>.
- 868 Haotian Zhang, Pengchuan Zhang, Xiaowei Hu, Yen-Chun Chen, Liunian Harold Li, Xiyang  
 869 Dai, Lijuan Wang, Lu Yuan, Jenq-Neng Hwang, and Jianfeng Gao. Glipv2: Uni-  
 870 fying localization and vision-language understanding. In Sanmi Koyejo, S. Mohamed,  
 871 A. Agarwal, Danielle Belgrave, K. Cho, and A. Oh (eds.), *Advances in Neural In-  
 872 formation Processing Systems 35: Annual Conference on Neural Information Process-  
 873 ing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9,  
 874 2022*, 2022. URL [http://papers.nips.cc/paper\\_files/paper/2022/hash/ea370419760b421ce12e3082eb2ae1a8-Abstract-Conference.html](http://papers.nips.cc/paper_files/paper/2022/hash/ea370419760b421ce12e3082eb2ae1a8-Abstract-Conference.html).
- 875 Xiao Zhang and Michael Maire. Self-supervised visual representation learning from hierarchical  
 876 grouping. In Hugo Larochelle, Marc'Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan,  
 877 and Hsuan-Tien Lin (eds.), *Advances in Neural Information Processing Systems 33: Annual  
 878 Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12,  
 879 2020, virtual*, 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/c1502ae5a4d514baec129f72948c266e-Abstract.html>.
- 880 Tiancheng Zhao, Tianqi Zhang, Mingwei Zhu, Haozhan Shen, Kyusong Lee, Xiaopeng Lu, and  
 881 Jianwei Yin. Vi-checklist: Evaluating pre-trained vision-language models with objects, attributes  
 882 and relations. *CoRR*, abs/2207.00221, 2022. doi: 10.48550/ARXIV.2207.00221. URL <https://doi.org/10.48550/arXiv.2207.00221>.
- 883 Yiwu Zhong, Jianwei Yang, Pengchuan Zhang, Chunyuan Li, Noel Codella, Liunian Harold Li,  
 884 Luowei Zhou, Xiyang Dai, Lu Yuan, Yin Li, and Jianfeng Gao. Regionclip: Region-based  
 885 language-image pretraining. In *IEEE/CVF Conference on Computer Vision and Pattern Recog-  
 886 nition, CVPR 2022, New Orleans, LA, USA, June 18-24, 2022*, pp. 16772–16782. IEEE, 2022.  
 887 doi: 10.1109/CVPR52688.2022.01629. URL <https://doi.org/10.1109/CVPR52688.2022.01629>.
- 888  
 889  
 890  
 891  
 892  
 893  
 894  
 895  
 896  
 897  
 898  
 899  
 900  
 901  
 902  
 903  
 904  
 905  
 906  
 907  
 908  
 909  
 910  
 911  
 912  
 913  
 914  
 915  
 916  
 917

## 918 A GENERATING BOX INFORMATION 919

920 We derive box information using an image grounding pipeline similar to Peng et al. (2023). Given  
921 an image-caption pair, noun entities are initially extracted from the caption into a list using spaCy  
922 (Honnibal et al., 2020). To minimize noise, we remove abstract nouns such as  $\{life, humor, love, \dots\}$   
923 from the list. We then predict the bounding boxes of the extracted entities within the image using the  
924 pre-trained grounding model GLIP (Li et al., 2022; Zhang et al., 2022). We exclude boxes with sizes  
925 lower than  $32 \times 32$ . We also threshold the predictions to at least 0.65 CLIP confidence score for the  
926 generated bounding box with the corresponding noun entity. Image-caption pairs for which no boxes  
927 could be generated or retained while filtering, are dropped. Further, referring expressions for noun  
928 chunks taken from the dependency tree of the caption using spaCy, are also included as text boxes.  
929 This increases the robustness of the stem towards linguistic complexities. **A few samples from the**  
930 **GRIT dataset are visualized in Fig. 7.**

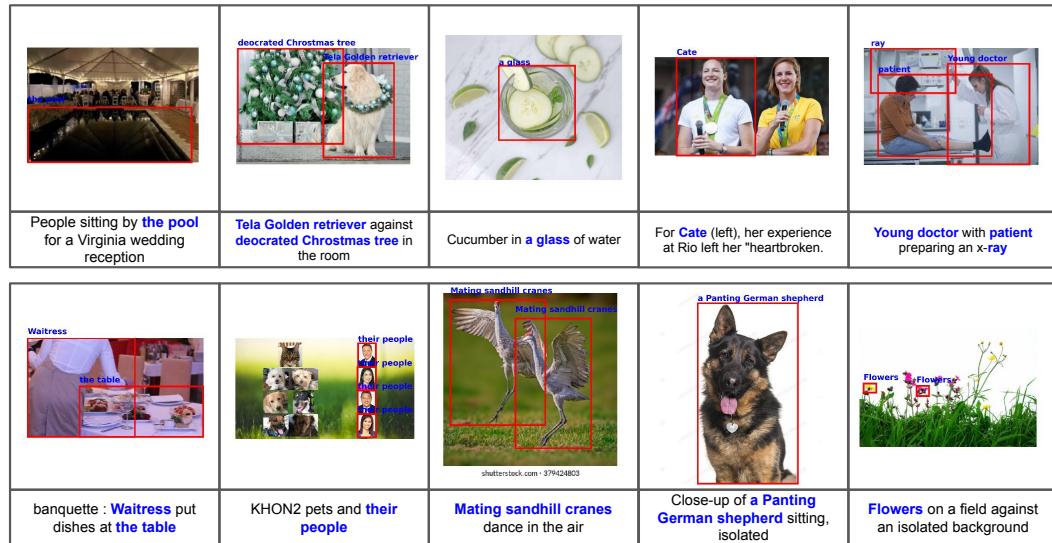
## 931 B IMPLEMENTATION DETAILS 932

933 **Model architecture** We use a similar setup as Desai et al. (2023),  
934 where the language encoder is the same one used by the original  
935 CLIP (Radford et al., 2021) consisting of a 12-layer Transformer  
936 architecture (Vaswani et al., 2017) with a width of 512 dimensions.  
937 The maximum input token size is set to 77 with a vocab size of 49408.  
938 For the vision encoder, we use the small and base Vision Transformer  
939 (Dosovitskiy et al., 2021; Chen et al., 2021; Touvron et al., 2021)  
940 backbone using a patch size of 16. The images are resized using  
941 border padding and random cropping (with scale  $[0.5, 1.0]$ ) to  $224 \times$   
942  $224$ , which results in an input sequence size of 196. A fixed set of 2-  
943 D sine-cosine position embeddings is included in the input sequence  
944 to instill a positional inductive bias.

945 **Initializing Lorentz model and Loss** We train HyCoCLIP with a  
946 fixed curvature value of the Lorentz model on the grounded CC3M  
947 dataset for 40k steps. To find the optimal setting, we additionally  
948

949 **Table 7: Training HyCo-  
950 CLIP on various settings of  
951  $\kappa$  on grounded CC3M dataset.  
952  $dnc$  denotes did not converge.**

$\kappa$	ImageNet
Fixed param.	
1.0	15.2
0.8	<i>dnc</i>
0.6	14.7
0.3	15.1
0.1	16.0
Learnable param.	
<b>1.0</b>	<b>16.7</b>



953 **Figure 7: Samples from the GRIT dataset.** The GRIT dataset (Peng et al., 2023) contains 20.5  
954 million grounded vision-language pairs of which 10 random samples are visualized in this illustration.  
955 The grounded information includes noun terms or their referring expressions (highlighted in blue)  
956 and the corresponding bounding boxes (in red within the image).

train by keeping the curvature a learnable parameter with an initial value of  $\kappa = 1.0$  and clamped within  $[0.1, 10.0]$ . As shown in Table 7, keeping the curvature a learnable parameter yields the best performance. Similar to Desai et al. (2023), we scale our batch of vectors before projecting it to the hyperboloid using learnable scalars  $c_{img}$  and  $c_{txt}$ , respectively, in both image and text modes. These scalars are initialized with a value of  $c_{img} = c_{txt} = 1/\sqrt{512}$ . The adaptive softmax temperature of the contrastive loss is initialized with  $\tau = 0.07$  and clipped at 0.01. All of these scalar values are learned in the logarithmic space.

In the hCE loss (Equations 10,11), we set separate values of the  $\eta$  parameter for inter-modality entailments  $\eta_{inter} = 0.7$  and intra-modality entailments  $\eta_{intra} = 1.2$  through a hyperparameter search while pre-training on the CC3M dataset for 75k steps and evaluating on ImageNet zero-shot image classification (cf. Table 8). Intuitively, this is because embeddings of images and text exist in different regions of the space, making it easier for text to entail the corresponding image as texts are nearer the origin and have a wider aperture  $\omega$  (cf. Eq. 7). Hence, we make the loss stricter by reducing the aperture of text embeddings. Similarly, the intra-modal box representations are closer in their corresponding spaces. Accordingly, we increase the aperture of the box regions to relax the entailment loss. In the final hC loss, we set the weight for hCE loss  $\gamma = 0.1$ .

**Optimizer and Hardware** We train our models on 4 A100 GPUs for 500k steps using a batch size of 768 on an internal cluster. Similar to Desai et al. (2023), we use the AdamW optimizer (Loshchilov & Hutter, 2019) with hyperparameters  $\beta_1 = 0.9$ ,  $\beta_2 = 0.98$  and weight decay 0.2 which is disabled for the learnable scalars. We use a cosine learning rate scheduler (Loshchilov & Hutter, 2017) with a maximum learning rate of  $5 \times 10^{-4}$  and a linear rate for the initial 4k steps.

## C METRICS FOR HIERARCHICAL CLASSIFICATION

This section provides more details on the metrics used for our hierarchical classification experiment. For a pair of predicted and true class  $(\hat{y}, y)$ , the Tree Induced Error (TIE) (Dekel et al., 2004) is the distance between  $\hat{y}$  and  $y$  in the graph (cf. Fig. 8a). This is defined as  $\sum_{e \in E(\hat{y}, y)} w_e$ , where  $E(i, j)$  is the set of edges with weights  $w_e$  along the path connecting nodes  $i$  and  $j$ . For the WordNet graph, we set  $w_e = 1$ . Similarly, the Lowest Common Ancestor (LCA) error is the distance to the deepest common node in the graph which is shared between the ancestors of  $\hat{y}$  and  $y$ .

For set-based measures, we define  $\hat{Y}_{anc}$  and  $Y_{anc}$  as the set of ancestor nodes of classes  $\hat{y}$  and  $y$  respectively (cf. Fig. 8b). Other relevant set-based hierarchical metrics such as, Jaccard Similarity  $J$ , and Hierarchical precision ( $P_H$ ) and recall ( $R_H$ ) (Kosmopoulos et al., 2015) are then given by

$$J = \frac{|\hat{Y}_{anc} \cap Y_{anc}|}{|\hat{Y}_{anc} \cup Y_{anc}|}, \quad P_H = \frac{|\hat{Y}_{anc} \cap Y_{anc}|}{|\hat{Y}_{anc}|}, \quad R_H = \frac{|\hat{Y}_{anc} \cap Y_{anc}|}{|Y_{anc}|}. \quad (13)$$

## D REGIONCLIP ON OTHER TASKS

In addition to object detection (cf. Sec. 3.2, **Zero-shot object detection**), we evaluate RegionCLIP on the other downstream tasks described in Sec. 3.2. We compare RegionCLIP, which uses a ResNet-50 backbone ( $\sim 25$ M parameters), to HyCoCLIP, which employs a ViT-S/16 backbone with  $\sim 22$ M parameters. As shown in Table 9, our method significantly outperforms RegionCLIP across all these tasks. This is expected, as RegionCLIP is primarily optimized for object detection. By contrast, HyCoCLIP reaches the level of performance of RegionCLIP on object detection despite being designed to learn hierarchical representation spaces.

## E SCENE UNDERSTANDING BENCHMARKS

In this section, we describe in detail the experiments of the compositional reasoning benchmarks used to evaluate our models in Sec. 3.2.

Table 8: **Hyperparameter search for  $\eta_{inter}$**  performed on grounded CC3M dataset.

$\eta_{inter}$	ImageNet
1.0	12.5
0.9	12.6
0.8	13.1
<b>0.7</b>	<b>13.4</b>
0.6	13.3
0.5	12.8

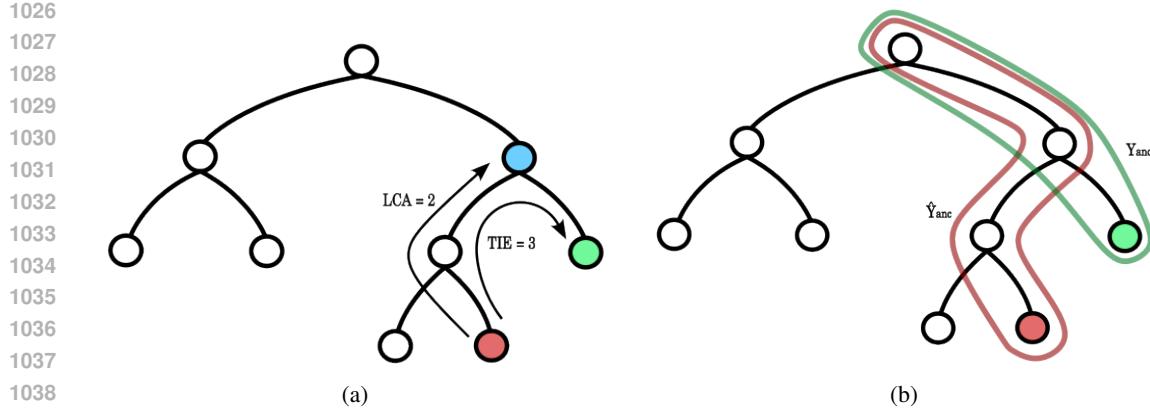


Figure 8: **Hierarchical classification metrics** when labels have a tree-like structure. Predicted label in **red** and true label in **green**. (a) Tree Induced Error (TIE) is the graph distance till the correct label while Least Common Ancestor (LCA) is the distance till the first common ancestor label as shown in **blue**. (b) shows ancestor label set of the predicted label ( $\hat{Y}_{anc}$ ) and the true label ( $Y_{anc}$ ).

Table 9: **RegionCLIP performance on downstream tasks.** The pre-training of RegionCLIP using boxes is primarily optimized for object detection. Thus displays poorer performance in comparison to HyCoCLIP which is designed to learn hierarchical representations from the boxes. \* s.u. denotes scene understanding.

Model	classification				retrieval		hierarchical metrics				s.u.*		
	ImageNet	CIFAR-100	SUN397	Food-101	Mean(16)	Text	Image	TIE ( $\downarrow$ )	LCA ( $\downarrow$ )	$J(\uparrow)$	$P_H(\uparrow)$	$R_H(\uparrow)$	VL-CO
	RegionCLIP	40.6	23.2	43.4	41.3	36.4	38.5	31.5	3.76	2.29	0.77	0.84	52.5
HyCoCLIP	<b>41.7</b>	<b>53.6</b>	<b>52.5</b>	<b>50.2</b>	<b>41.1</b>	<b>69.5</b>	<b>55.2</b>	<b>3.55</b>	<b>2.17</b>	<b>0.79</b>	<b>0.86</b>	<b>0.85</b>	<b>59.8</b>

## E.1 BENCHMARKS

**VL-Checklist-Object (VL-CO)** This benchmark (Zhao et al., 2022) modifies the caption in several aspects. An object term in the caption is replaced with a random noun phrase. The model results are categorized for different sizes and locations of the object within the image to check for invariance which are summarized as follows,

- **O-Small:** The object covers a small area within the image. Following Zhao et al. (2022), the threshold of the object area is set to below  $32 \times 32$ .
- **O-Medium:** The object covers a moderate area within the image. The threshold of the object area is set between  $32 \times 32$  and  $96 \times 96$ .
- **O-Large:** The object covers a large area within the image. Any object with an area greater than  $96 \times 96$  fits this category.
- **O-Center:** The object center lies within the center region of the image. If  $x$  is the half-length diagonal, and  $y$  is the distance between the center of the object and the center of the image, the object is considered to lie in the center region if  $\frac{y}{x} \leq \frac{1}{3}$ .
- **O-Margin:** The object center lies at the margin of the image. This is when  $\frac{y}{x} > \frac{2}{3}$ .
- **O-Mid:** The object center lies in between the center and margin region which is when  $\frac{1}{3} < \frac{y}{x} \leq \frac{2}{3}$ .

**VG-Attribution (VG-A)** This benchmark (Yüksekgün̄l et al., 2023) tests the capability of the model to correctly identify the attribute word associated with an object in a sentence in the context of an image. For instance, the model has to pick between “the **crouched cat** and the **open door**” and “the **open cat** and the **crouched door**”.

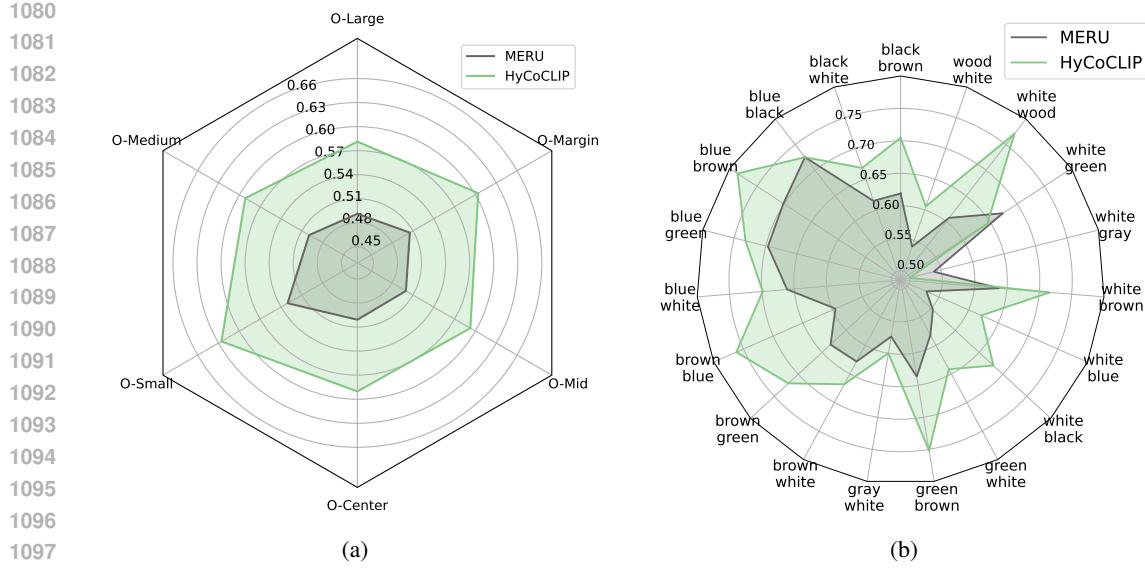


Figure 9: **Performance on scene understanding benchmarks** (a) model performance on VL-Checklist-Object (VL-CO) understanding experiments. HyCoCLIP performs best on object understanding tasks. (b) model performance on Visual Genome (VG) Attribution benchmark, accuracy values are plotted for the 19 most frequent attribute pairs in the experiment. HyCoCLIP gives the best results on most of these attribute pairs.

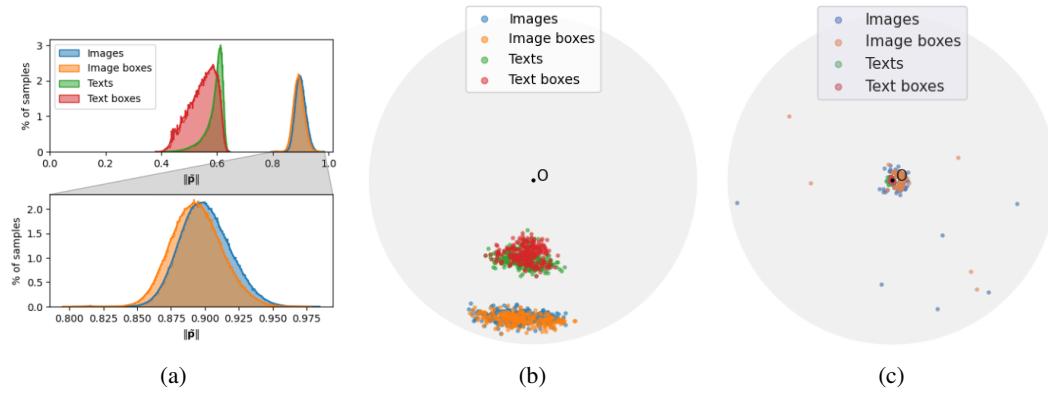


Figure 10: **Visualizing the learned hyperbolic space of MERU in lower dimensions** using samples from GRIT. (a) distribution of embedding distances from the origin, MERU embeds text data closer to the origin wrt the images but box samples don't show a smaller radius wrt full images/captions. On the right, (b) HoroPCA does not show local ordering and (c) CO-SNE visualizations of the latent space in  $\mathbb{L}^2$  are quite distorted.

## E.2 PERFORMANCE

We reported the mean performance for scene understanding benchmarks in Table 5. In addition to this, here we provide the results of HyCoCLIP compared against MERU on the individual categories of VL-CO and the top 19 most frequently occurring attribute pairs in the VG-A evaluation set. Fig. 9 highlights the significant improvements achieved by our method indicating better semantic scene understanding.

1134 F VISUALIZING REPRESENTATION SPACE - MERU  
 1135  
 1136

1137 Following Sec. 4, we additionally plot the learned hyperbolic space of MERU in lower dimensions in  
 1138 Fig. 10. We observe that the distributions of text embeddings and image embeddings are overlapped  
 1139 with corresponding box distributions from Fig. 10a. This is also apparent in the 2D plot when  
 1140 using HoroPCA in Fig. 10b. The embeddings for all modes seem to collapse to a small region with  
 1141 CO-SNE as seen in Fig. 10c.

1142  
 1143 G INTERPOLATION DETAILS  
 1144

1145 The logarithmic map is the inverse of the exponential map and, for  $\mathbf{p}, \mathbf{q} \in \mathbb{L}^n$ , is given by  
 1146

$$\log_{\mathbf{p}}^{\kappa}(\mathbf{q}) = \frac{\cosh^{-1}(-\kappa \langle \mathbf{p}, \mathbf{q} \rangle_{\mathbb{L}})}{\sqrt{(\kappa \langle \mathbf{p}, \mathbf{q} \rangle)^2 - 1}} (\mathbf{q} + \kappa \langle \mathbf{p}, \mathbf{q} \rangle_{\mathbb{L}} \mathbf{p}). \quad (14)$$

1147 Here, we will use it to interpolate between points in hyperbolic space (Desai et al., 2023).  
 1148

1149 For hyperbolic representations  $(g_I(I_S), g_I(I_T))$  of source image  $I_S$  and target image  $I_T$ , we take the  
 1150 logarithmic maps  $\log_0(g_I(I_S))$  and  $\log_0(g_I(I_T))$  and obtain a set of  $N$  equally spaced representations  
 1151 on the line joining these vectors given by,  
 1152

$$S_E^N = \left\{ p_i \in \mathcal{T}_p \mathbb{L}^n : p_i = (1 - t_i) \log_0(g_I(I_S)) + t_i \log_0(g_I(I_T)), t_i = \frac{i}{N}, i \in \{1, \dots, N\} \right\}. \quad (15)$$

1153 These are then mapped back to the hyperboloid using exponential mapping given by,  
 1154

$$S_H^N = \left\{ q_i \in \mathbb{L}^n : q_i = \exp_0(p_i), p_i \in S_E^N \right\}. \quad (16)$$

1155  
 1156 The closest representations are retrieved for all points in the set  $S_H^N$  using Lorentzian distance as  
 1157 the similarity metric from a collection of representations of images and texts. Further, we drop any  
 1158 duplicate representations retrieved. **For hyperbolic models, the origin (0) of the space is taken as**  
 1159 **the root node, whereas for Euclidean CLIP, the root node is taken as the centroid of all the training**  
 1160 **samples.**

1161  
 1162 H HIERARCHICAL IMAGE-TEXT MATCHING BENCHMARK  
 1163

1164 In this section, we evaluate our model and baselines  
 1165 on the hierarchical image-text matching benchmark  
 1166 using the HierarCaps (Alper & Averbuch-Elor, 2024)  
 1167 test set. Similar to the setup of Alper & Averbuch-  
 1168 Elor (2024) and Appendix G, we take 50 equally-  
 1169 spaced points between the root node and the nearest  
 1170 text embedding to the image embedding and calcu-  
 1171 late the precision (P) and recall (R) relative to the  
 1172 four hierarchically relevant ground-truth captions. In  
 1173 addition, we also report the order-aware metric  $\tau_d$ ,

1174 which shows if the captions are ordered correctly in the embedding space. From the results in  
 1175 Table 10, we find HyCoCLIP outperforms CLIP and MERU, further demonstrating the enhanced  
 1176 hierarchical understanding achieved with our method. Further, from Fig. 11 we empirically find  
 1177 HyCoCLIP performs better hierarchical image-text matching on HierarCaps.

Table 10: **Hierarchical text-image matching**  
 performed on HierarCaps benchmark.

Model	P	R	$\tau_d$
CLIP	0.13	0.29	0.83
MERU	0.12	0.39	0.84
HyCoCLIP	0.12	0.46	0.88

1188  
1189  
1190  
1191  
1192  
1193



HyCoCLIP	MERU	CLIP
remark	something	looking
out	image	legume
something	looking	opening
crop	photograph	variety
good	a picture of food .	crop
cooking	eukaryotic food	bud
savory	fresh food	a blinder that has some vegetables in it
vegetable	legume	lettuce
some vegetables	plant vegetable	A opened peeled banana half way eaten.
green veggie	some gross ass vegetables that look slimy AF	
edible vegetable	edible vegetable	
broccoli	herb	
a hand holding a vegetable	fried vegetables	
plate of broccoli	a blinder that has some vegetables in it	
a piece of broccoli	a hand holding a vegetable	
fresh broccoli with a salad	A raw piece of broccoli with something growing from it.	
finger holding a broccoli	Two fingers holding a small piece of broccoli	
The broccoli platter in the dish is mostly eaten up.	finger holding a broccoli	
A worm sits on top of a piece of broccoli.		

1201  
1202  
1203  
1204  
1205  
1206  
1207  
1208  
1209  
1210  
1211  
1212  
1213  
1214



HyCoCLIP	MERU	CLIP
remark	something	looking
out	picture	behind
something	looking	soft
miles	good	joes
little	people	stuffed
stuffed	baking food	teddy
some stuff	edible snack	stuffed bear
junk	sugary sweet	stuffed animal in a food
tiny room	affection	teddy bears
toys	two	a stuffed bear with others
teddy	fried fried	close up of stuffed animals .
stuffed bear	natural fiber	a couple of brown teddy bears
a little stuffed bear	stuffed	the teddy bears lying down on a quilt
stuffed bears	teddy	
a stuffed teddy bear sitting	a little stuffed bear	Three stuffed bears hugging and sitting on a blue pillow
a pair of stuffed teddy bears	stuffed teddy bear	
a stuffed bear with others	stuffed teddy bears	
Two stuffed teddy bears resting next to each other	stuffed bears	
a stuffed bear and a teddy bear	a pair of stuffed teddy bears	
	a stuffed bear with others	
Three teddy bears, each a different color, snuggling together.	THIS IS A CLOSE UP PICTURE OF A STUFFED BEAR AND MONKEY	
Three teddy bears laying in bed under the covers.	a couple of brown teddy bears	
Two teddy bears are laying side by side on a quilt.	teddy bears resting next to each other	
There are two stuffed bears and one of them is wearing a shirt.	Three teddy bears, each a different color, snuggling together.	
Two stuffed animals sit at a table with honey.	Two stuffed animals sit at a table with honey.	

1215  
1216  
1217  
1218  
1219  
1220  
1221  
1222  
1223  
1224  
1225  
1226  
1227  
1228  
1229  
1230  
1231  
1232  
1233  
1234  
1235  
1236  
1237  
1238  
1239  
1240  
1241

HyCoCLIP	MERU	CLIP
remark	something	looking
out	image	living
something	living thing	legume
instrumental	guy	serving
legume	a picture of food .	curmudgeon
serving	savory food	food
good	The asian dish is filled with several different ingredients.	food served
cooking		hot meal
dish	a dish made from food on a plate	a plate with cooked food
dinner	a plate	person & dish with fajita on plate .
fried	meal	
cooked food	a plate of food .	a plate with cooked rice , vegetables and chicken
rice	the meal on a plate	
hot meal	a plate of food in a cooking .	
chicken meal	a plate with cooked food	
the dish on a plate	an image of a plate of food with meat and veggies	
the meal on a plate		
a meal with some ingredients	person & dish with fajita on plate .	
a plate with cooked food and meat	A plate of vegetables, chicken, and white rice.	
a plate with cooked food	A plate of chicken and vegetables sits next to a bowl of rice.	
a plate with cooked rice , vegetables and chicken		
A dinner plate that has white steamed rice with stir fry vegetables and chicken.		
A plate of vegetables, chicken, and white rice.		
A good looking dish of food is arranged on a plate		
A plate that has food that is sitting on the table.		

HyCoCLIP	MERU	CLIP
remark	something	looking
bef	picture	soft
ohone	looking	pols
little	plural	lying
frizbee	domestic animal	furry
two	dilute	his fright
booie	cat	sleeping
lying	kitty	cats
furry	her cat	There are pets such as dogs that are sleeping in the bef
kitty	kitten	
feline	cub	pets sleeping in the bef
cat laying down	felina	sleeping cats
kitty sitting	pets sleeping in the bef	cats sleeping
cat lying down	sleeping cat	cats sleeping with a remote
sleeping cat		the cats rest on a comforter
sleeping cats	kitty sleeping	
cats resting	sleeping cats	Two cats rest on each other to take a nap.
cats sleeping		
cat resting on each other	the cats are lying down	Two cats lay together on a blanket.
Two cats laying together on a bed or floor	cats sleeping with a remote	A couple of cats laying on top of a pink blanket.
there are two cats that are laying with each other	the cats rest on a comforter	
Two cats lay together on a blanket.		
A couple of cats laying on top of a pink blanket.	two cats lying near two remotes on a purple sheets	
A young cat on a mat with a flip flop shoe.	A couple of cats laying on top of a pink blanket.	
two cats laying near two remotes on a purple sheets		

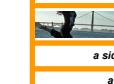
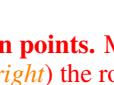
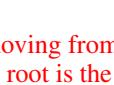
Figure 11: Hierarchical image-text matching qualitative results on HierarCaps test set.

1242 **I MORE INTERPOLATION EXAMPLES**

1243

1244

1245

SOURCE IMAGE					
HyCoCLIP	MERU	CLIP	HyCoCLIP	MERU	CLIP
					
<i>a bridge of old and new Asia</i>			<i>A typical picture of New York City</i>	<i>the New York City skyline</i>	<i>New Yorkers</i>
<i>a photograph of an indian landmark</i>			<i>skyline</i>	<i>the skyline of New York</i>	<i>a high place</i>
			<i>the city</i>	<i>A typical picture of New York City</i>	<i>this photo</i>
			<i>middle</i>		<i>the pic</i>
ROOT					
TARGET IMAGE					
SOURCE IMAGE					
					
HyCoCLIP	MERU	CLIP	HyCoCLIP	MERU	CLIP
					
<i>A kite is being flown next to a suspension bridge, and the boat and the person flying the kite are visible in the foreground</i>	<i>harbor</i>		<i>A kite is being flown next to a suspension bridge, and the boat and the person flying the kite are visible in the foreground</i>	<i>harbor</i>	
<i>A person does a flip by two sailboats and a lake</i>	<i>a fish dinner</i>		<i>a boat sail past</i>		
<i>a fish dinner</i>			<i>the waterside</i>		
<i>A woman prepares a meal with fish, eggs, broccoli and celery in front of a blue curtain</i>			<i>the harbor</i>		
			<i>a scenic view</i>		
			<i>the photo</i>		
			<i>" Just</i>		
ROOT					
TARGET IMAGE					

1293 **Figure 12: Interpolation between points.** Multimodal retrieval results when moving from (*top*) an image to (*left*) another image or (*right*) the root. For HyCoCLIP and MERU, the root is the origin of the space, whereas it is the centroid of training sample representations for CLIP.

1294

1295

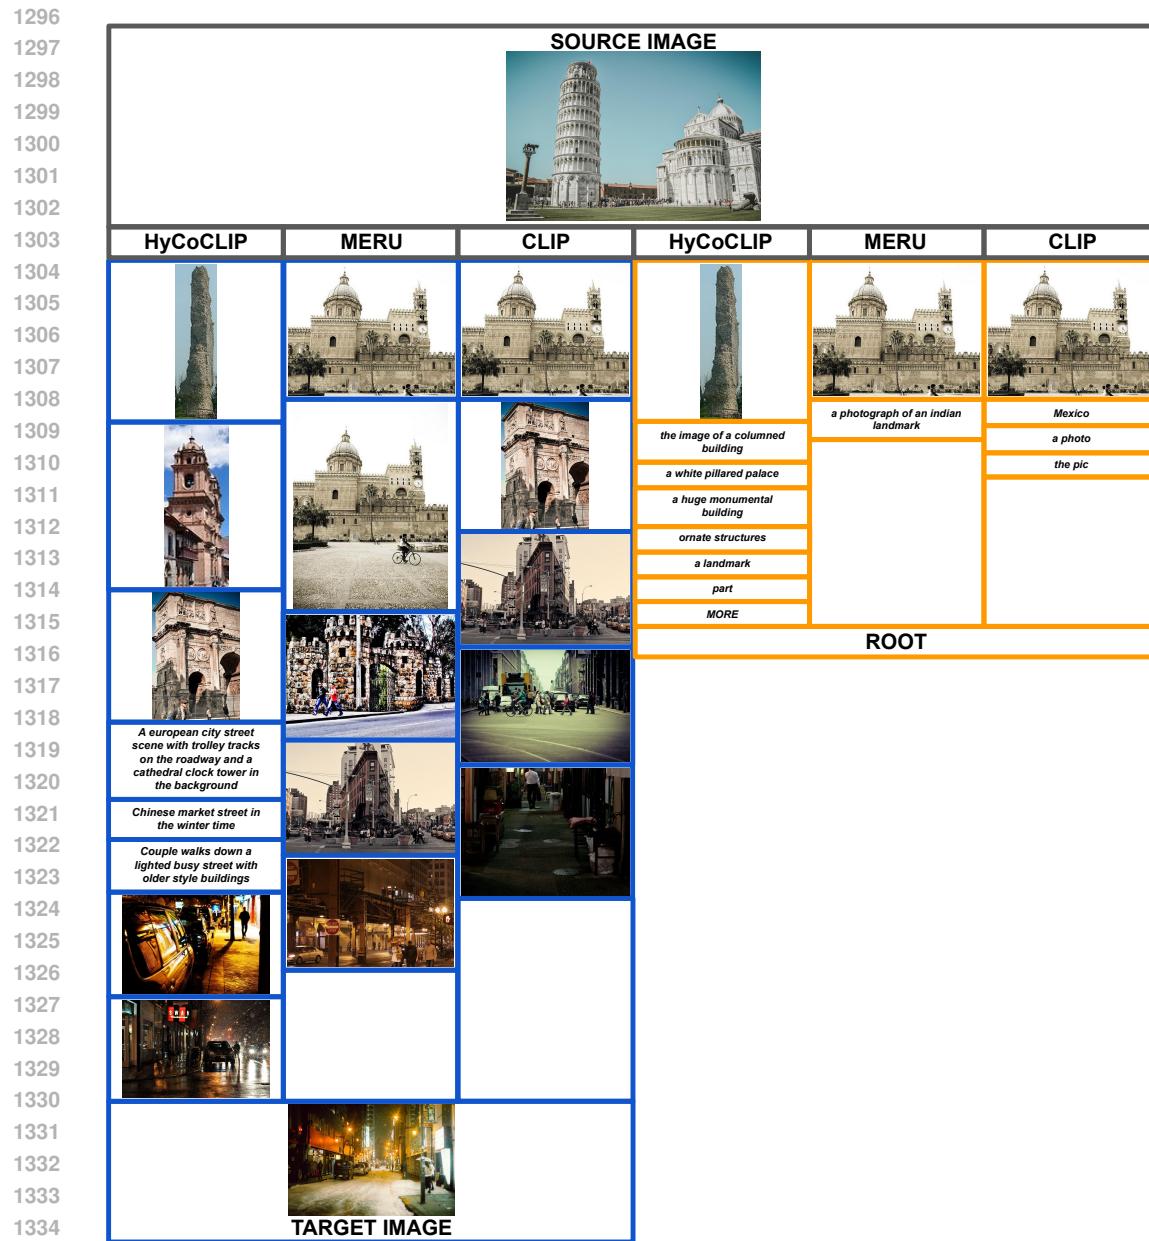


Figure 13: **Interpolation between points.** Multimodal retrieval results when moving from (*top*) an image to (*left*) another image or (*right*) the root. For HyCoCLIP and MERU, the root is the origin of the space, whereas it is the centroid of training sample representations for CLIP.

1350  
1351  
1352  
1353  
1354  
1355  
1356  
1357  
1358  
1359  
1360  
1361  
1362  
1363  
1364  
1365  
1366  
1367  
1368  
1369  
1370  
1371  
1372  
1373  
1374  
1375  
1376  
1377  
1378  
1379  
1380  
1381  
1382  
1383  
1384  
1385  
1386  
1387  
1388  
1389  
1390  
1391  
1392  
1393  
1394  
1395  
1396  
1397  
1398  
1399  
1400  
1401  
1402  
1403

SOURCE IMAGE					
HyCoCLIP	MERU	CLIP	HyCoCLIP	MERU	CLIP
some of the goods	numerous items		A bunch of items laid out across the ground	wood items	
other wooden items	various sunglasses		an assortment of paints	props	
red sea-saws	a bunch of sunglasses		a collection of stuff	bikinis	
	DX		assorted other items		
	a sea port		several other items		
	crystal blue waters		other items		
	A blue restaurant that sits in the water		a stuff		
			MORE		
			ROOT		
TARGET IMAGE					
SOURCE IMAGE					
HyCoCLIP	MERU	CLIP	HyCoCLIP	MERU	CLIP
many visible vapour trails			many visible vapour trails		
An adult riding a bike on a beach with many visible vapour trails in the sky			a light show		
A bicyclist becomes airborne among dirt hills at night			part		
				many visible vapour trails	
	many visible vapour trails			electricity	
	An adult riding a bike on a beach with many visible vapour trails in the sky			power	
				A picture of an artwork	
				this picture	
				an artwork	
				a pic	
				ROOT	
TARGET IMAGE					

Figure 14: **Interpolation between points.** Multimodal retrieval results when moving from (*top*) an image to (*left*) another image or (*right*) the root. For HyCoCLIP and MERU, the root is the origin of the space, whereas it is the centroid of training sample representations for CLIP.

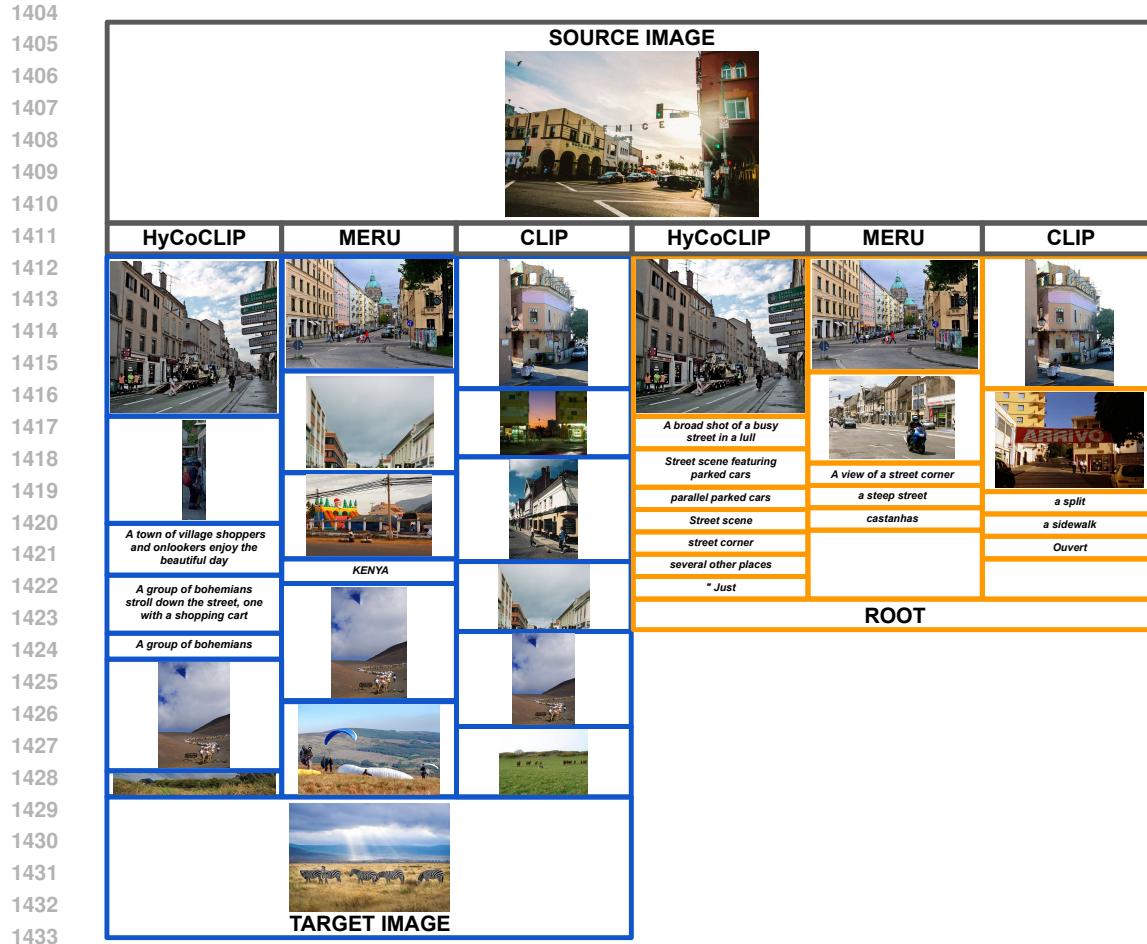


Figure 15: **Interpolation between points.** Multimodal retrieval results when moving from (*top*) an image to (*left*) another image or (*right*) the root. For HyCoCLIP and MERU, the root is the origin of the space, whereas it is the centroid of training sample representations for CLIP.

1458						
1459	<b>HyCoCLIP</b>	<b>MERU</b>	<b>CLIP</b>	<b>HyCoCLIP</b>	<b>MERU</b>	<b>CLIP</b>
1460						
1461	<i>a brightly lit railway passage</i>			<i>A photo of a train platform on a chilly day</i>	<i>A person crossing a bridge with train tracks</i>	<i>Train</i>
1462	<i>The people are walking by the train tracks</i>			<i>a brightly lit railway passage</i>	<i>a train line</i>	<i>railroads</i>
1463	<i>Two boys dressed in white capes running on a path in a field</i>	<i>the rails of a set</i>		<i>rail tracks</i>	<i>rails</i>	<i>Train</i>
1464		<i>rails</i>		<i>a train line</i>	<i>sucker</i>	<i>a sidewalk</i>
1465		<i>People trackside</i>		<i>station</i>		<i>a pic</i>
1466				<i>the picture</i>		<i>the pic</i>
1467				<i>part</i>		
1468				<i>middle</i>		
1469						
1470						
1471						
1472						
1473						
1474						
1475						
1476						
1477						
1478						
1479						
1480						
1481						
1482						
1483						
1484						
1485						
1486						
1487	<b>ROOT</b>					
1488	Figure 16: <b>Interpolation between points.</b> Multimodal retrieval results when moving from ( <i>top</i> ) an image to ( <i>left</i> ) another image or ( <i>right</i> ) the root. For HyCoCLIP and MERU, the root is the origin of the space, whereas it is the centroid of training sample representations for CLIP.					
1489						
1490						
1491						
1492						
1493						
1494						
1495						
1496						
1497						
1498						
1499						
1500						
1501						
1502						
1503						
1504						
1505						
1506						
1507						
1508						
1509						
1510						
1511						